Course Announcements

Due Sunday (11:59 PM):

- D7
- Q8
- Weekly Project Survey (*optional*)

Notes:

- A3 and D6 grades posted
- Data Checkpoint grading underway (discuss)
- update: A4 now available; due *Wed of week 10* (3/13)
    - Note: “Validate” will fail; use Kernel > Restart & Run All instead

Machine Learning I

Sid Joshi
UC San Diego

Department of Cognitive Science
s1joshi@ucsd.edu

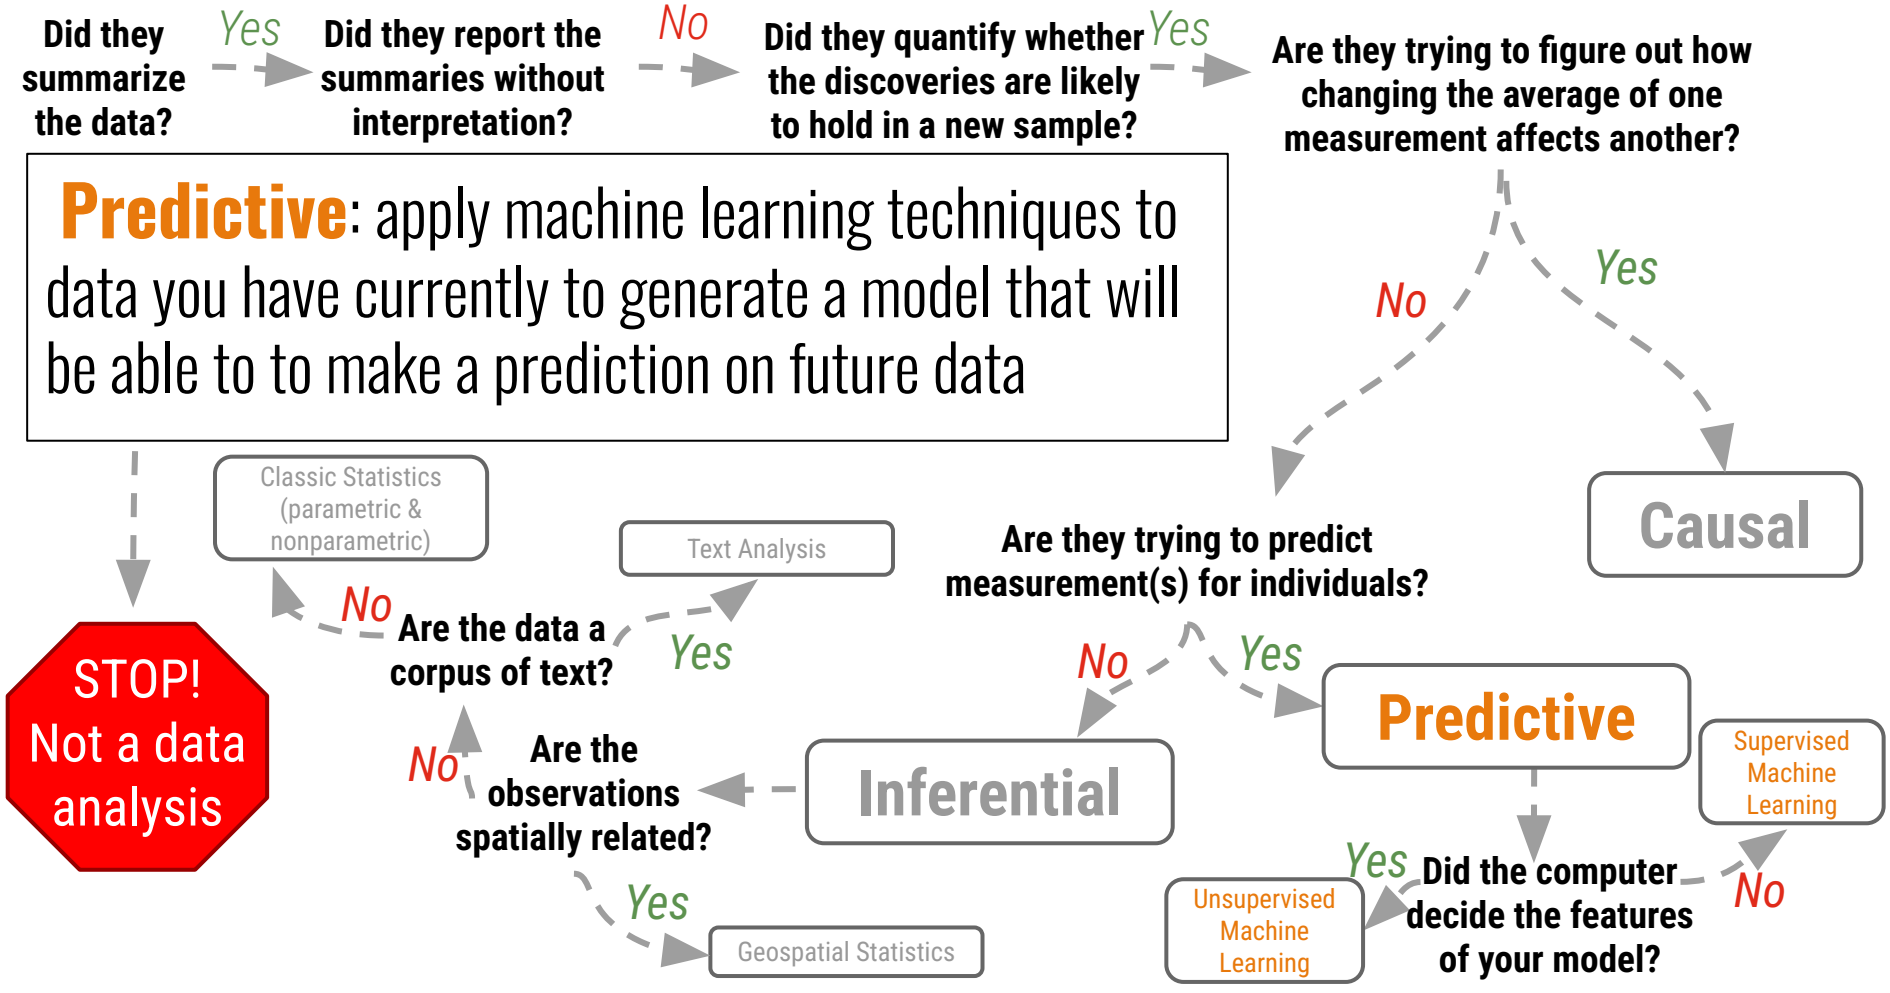
Did they summarize the data?
Yes
Did they report the summaries without interpretation?
No
Did they quantify whether the discoveries are likely to hold in a new sample?
Yes
Are they trying to figure out how changing the average of one measurement affects another?
Predictive: apply machine learning techniques to data you have currently to generate a model that will be able to to make a prediction on future data
No
Yes
Causal
Are they trying to predict measurement(s) for individuals?
No
Yes
Predictive
Supervised Machine Learning
Did the computer decide the features of your model?
Yes
Unsupervised Machine Learning
No
Inferential
Are the observations spatially related?
No
Yes
Geospatial Statistics
Are the data a corpus of text?
No
Yes
Text Analysis
Classic Statistics (parametric & nonparametric)
STOP! Not a data analysis

predictive analysis uses data you have now to make predictions in the future

machine learning approaches are used for predictive analysis!

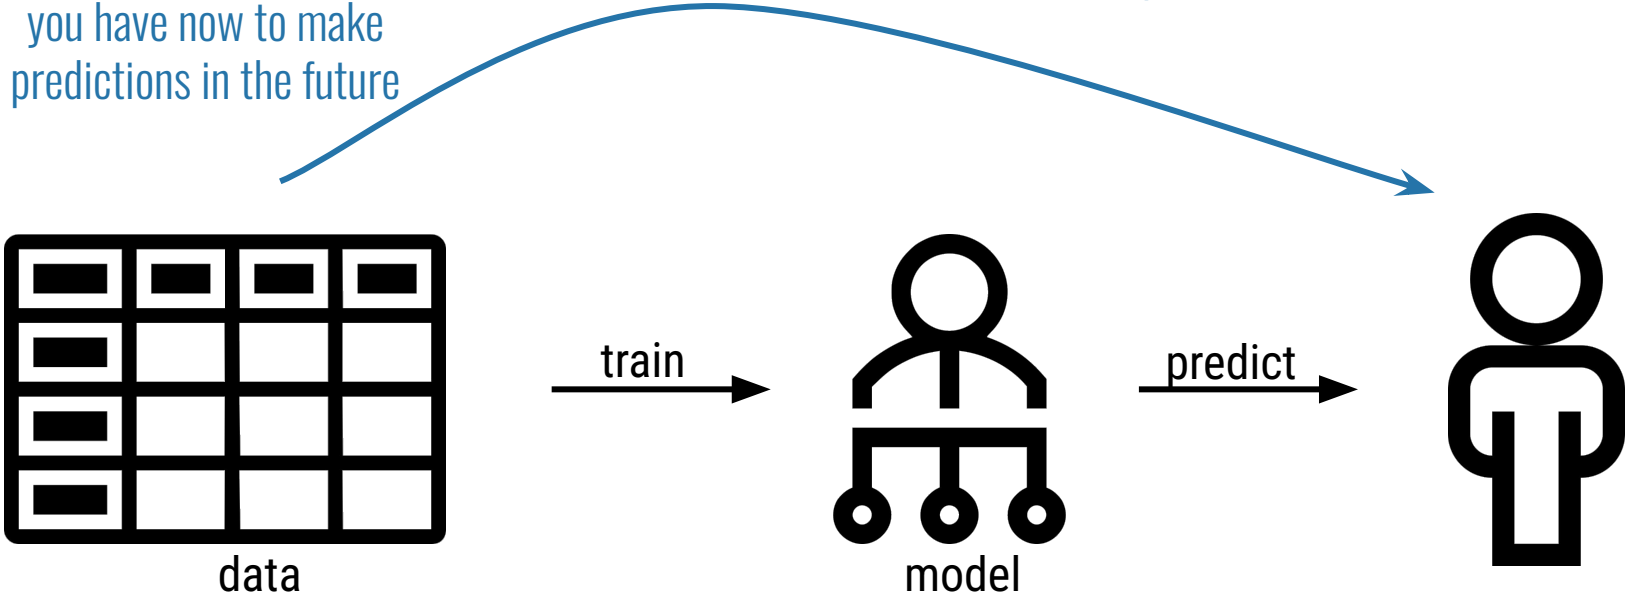

What is machine learning?

"Machine learning is the science of getting computers to act without being explicitly programmed"

- Andrew Ng, Stanford, ex-Google, chief scientist at Baidu, Coursera founder, Stanford Adjunct Faculty

A Shift in Programming Paradigm

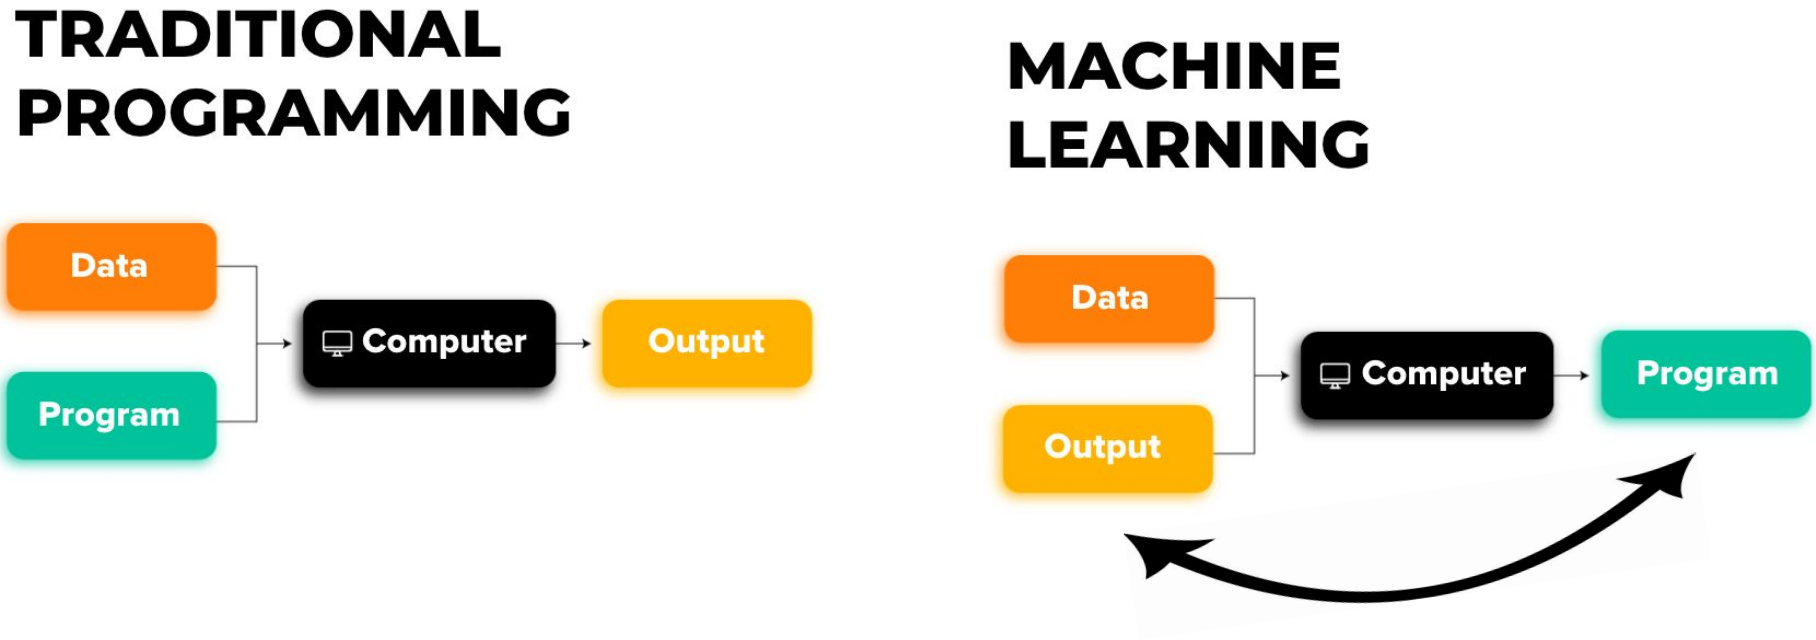

https://www.avenga.com/magazine/machine-learning-programming/

- **Problem:** Detecting whether credit card charges are fraudulent.
- **Data science question:** Can we use the time of the charge, the location of the charge, and the price of the charge to predict whether that charge is fraudulent or not?
- **Type of analysis:** Predictive analysis

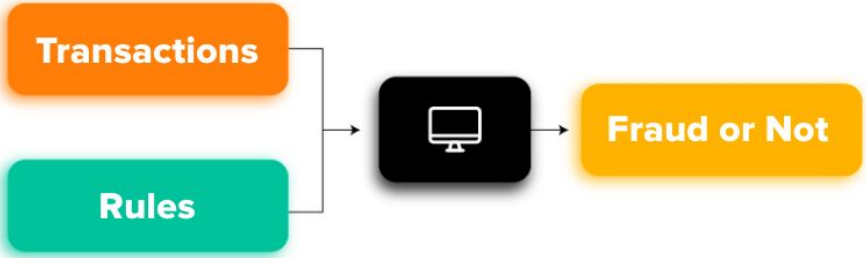

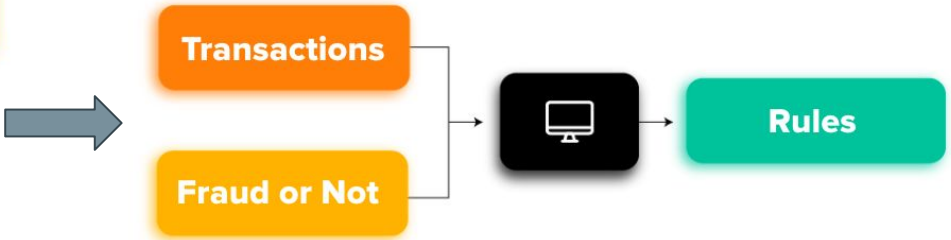

Rule1. Claim time - Submit time < 1 h
Rule2. Agreement review time > 5 m
Rule3. ...

Prediction Questions

Which of these questions is most appropriate for machine learning?

A How common is watching Sesame Street in the US?

B What is the effect of watching Sesame Street on children's brains?

C What is the relationship between early childhood educational programming and success in elementary school?

D Can we use information about one's early childhood to predict their success in elementary school?

E How does Sesame Street cause an increase in educational attainment?

Machine Learning Generalizations

The goal of modeling is to simplify, not replicate

Basic Steps to Prediction

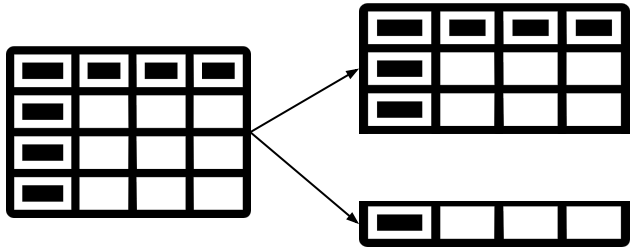

data
partitioning

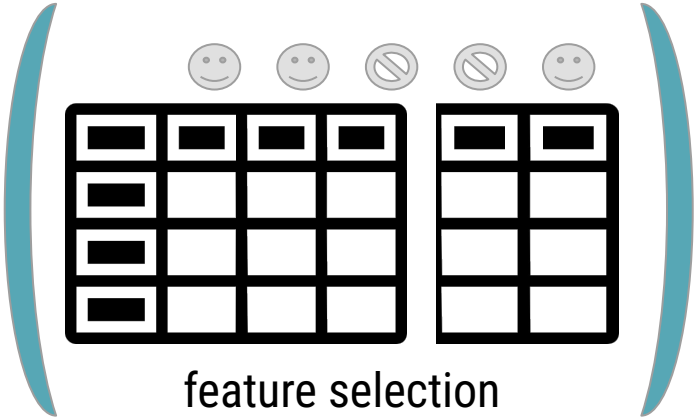

feature selection

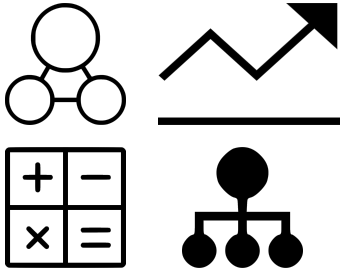

model selection

model assessment

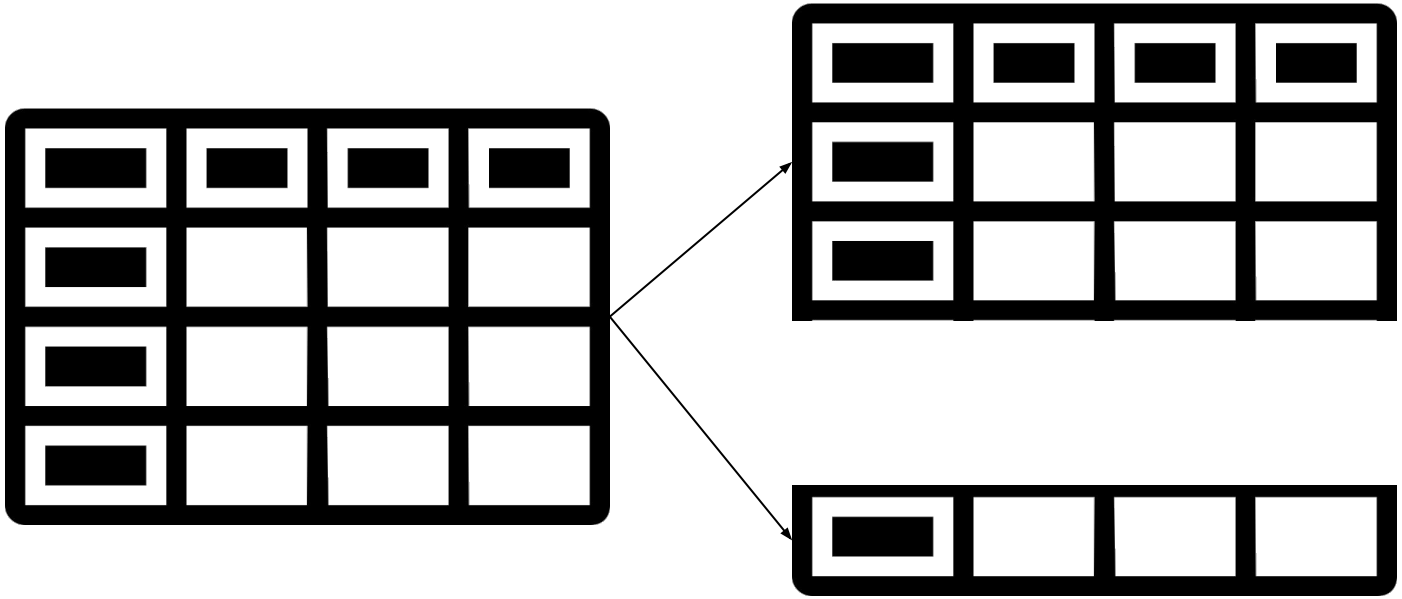

data partitioning

the data used to build your predictive model

Data set used to assess if prediction model is generalizable; can be held out subset or new data entirely

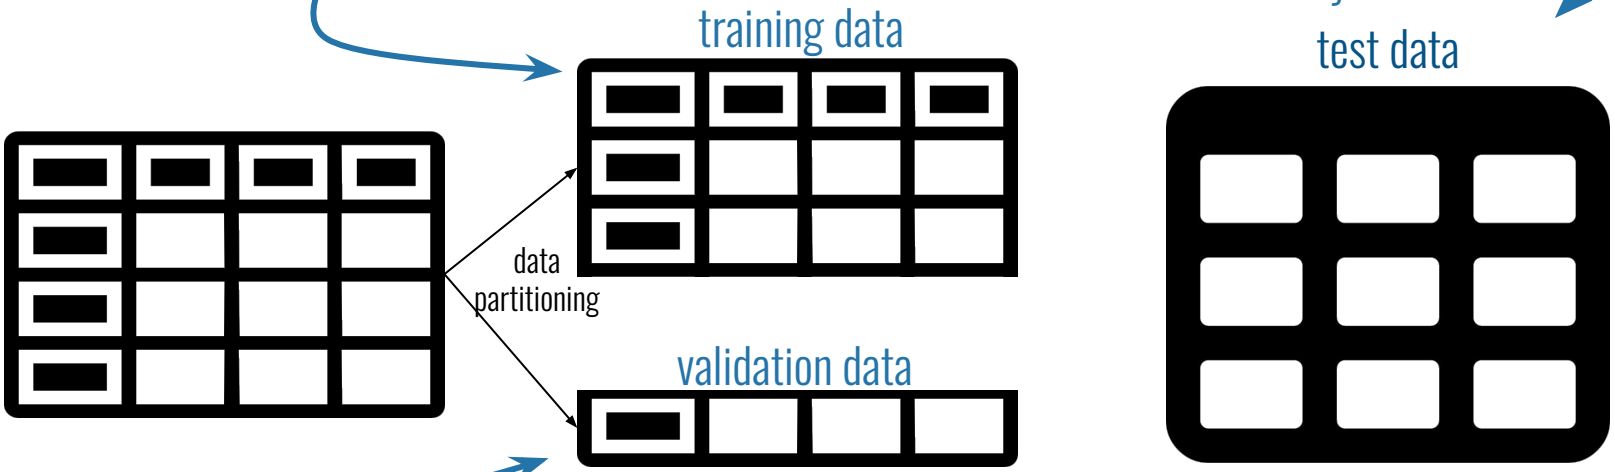

Data from original dataset that was held out and not used in training the model ; helpful in fine-tuning prediction accuracy

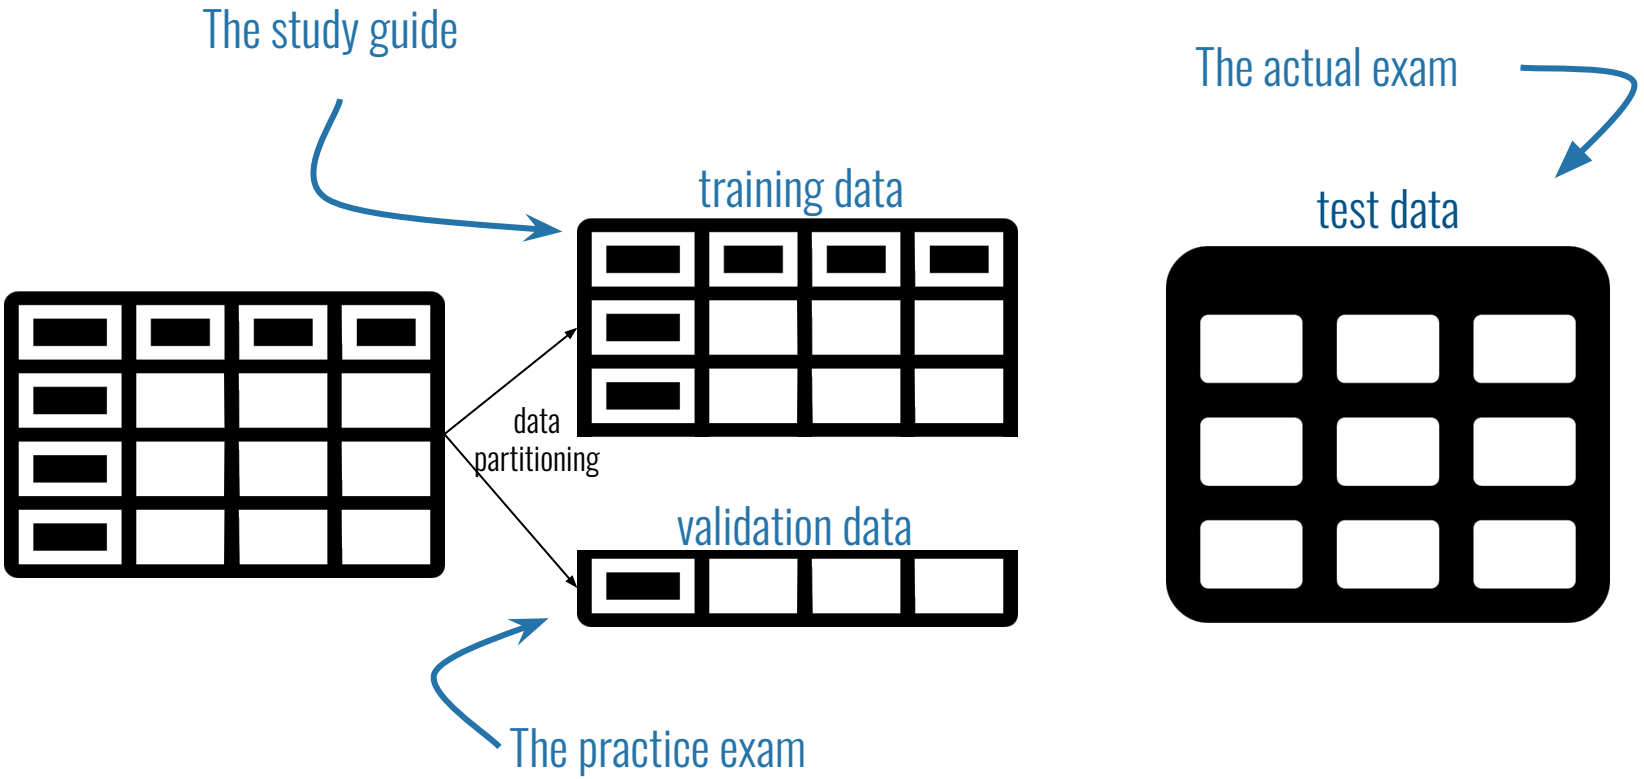
The study guide
The actual exam
training data
test data
data
partitioning
validation data
The practice exam

Data Partitioning

What portion of the data are typically used for generating the model?

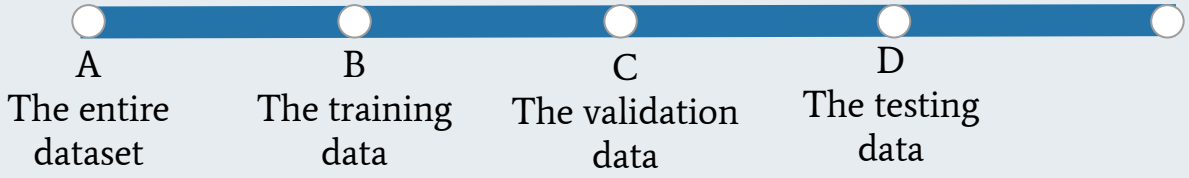

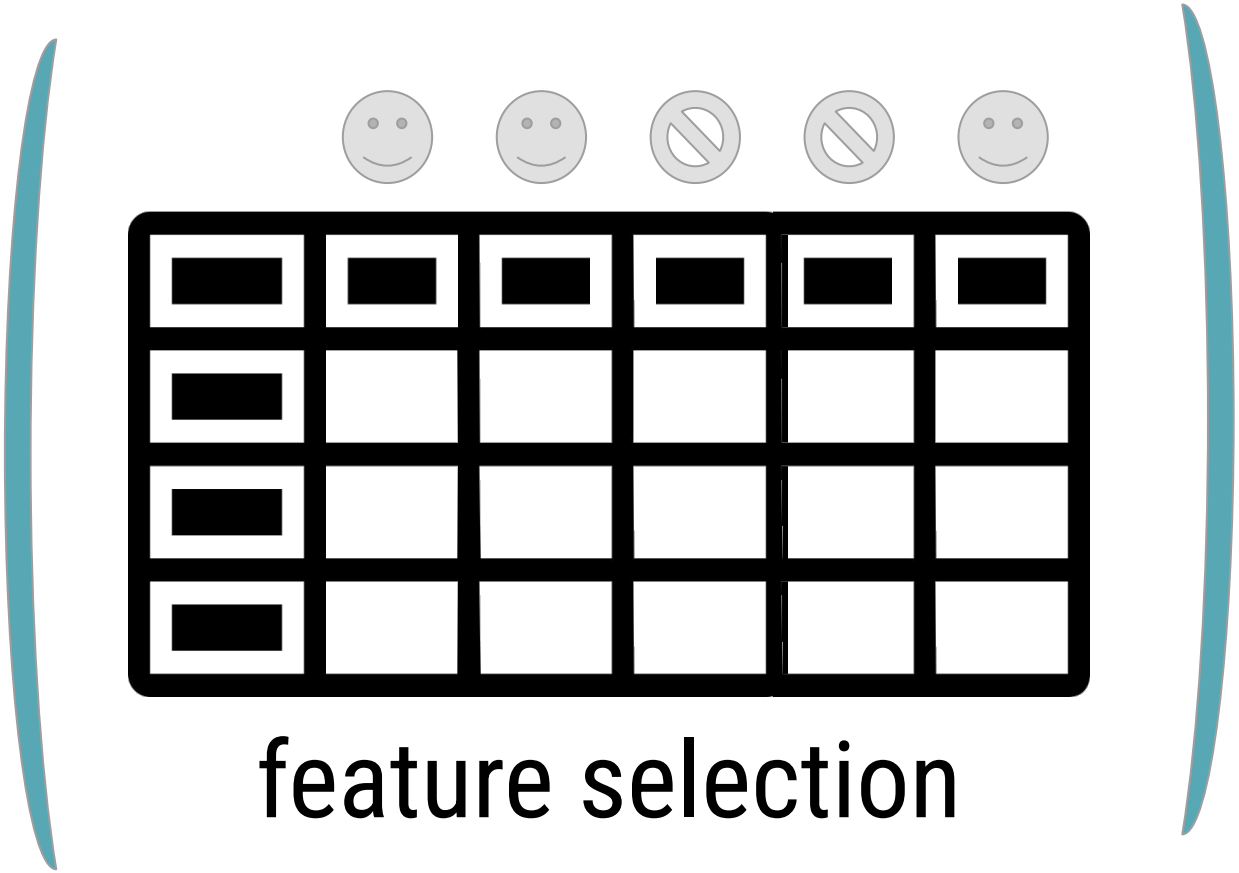
feature selection

elephant height data are likely
not predictive of US elections

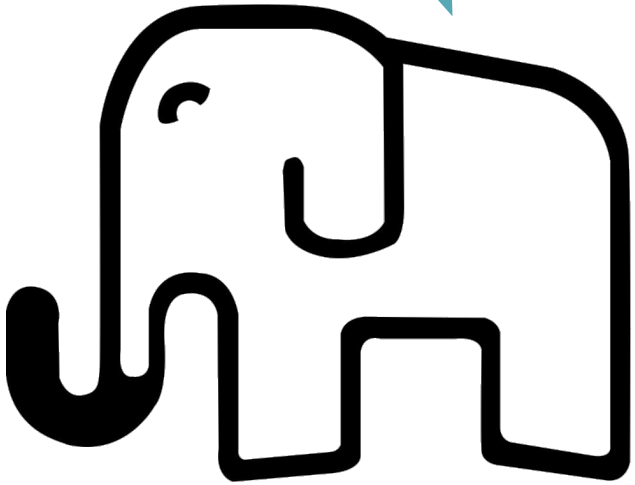

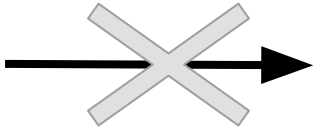

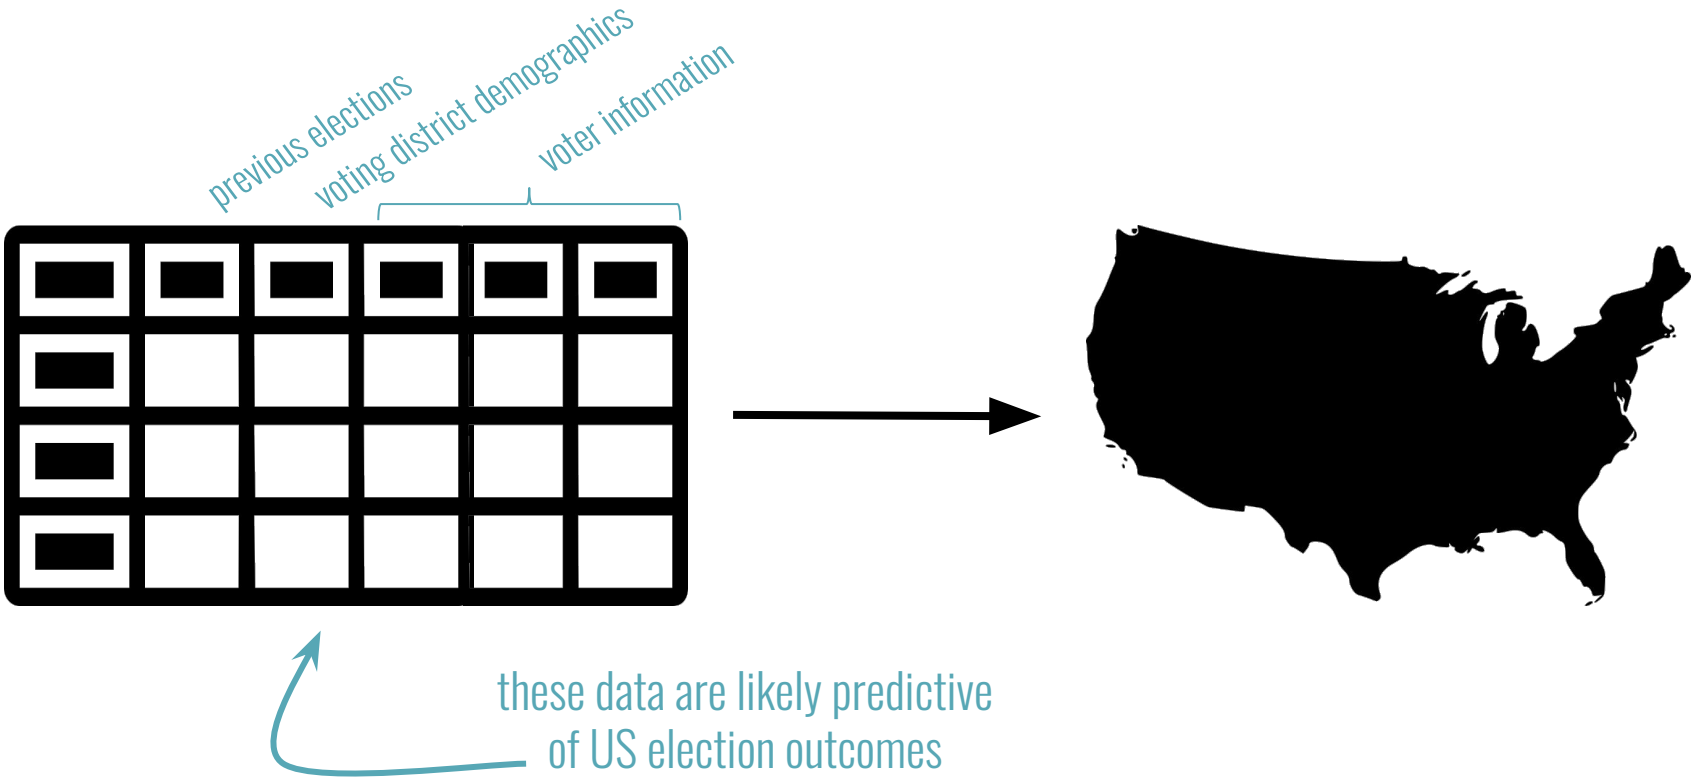
previous elections
voting district demographics
voter information
these data are likely predictive
of US election outcomes

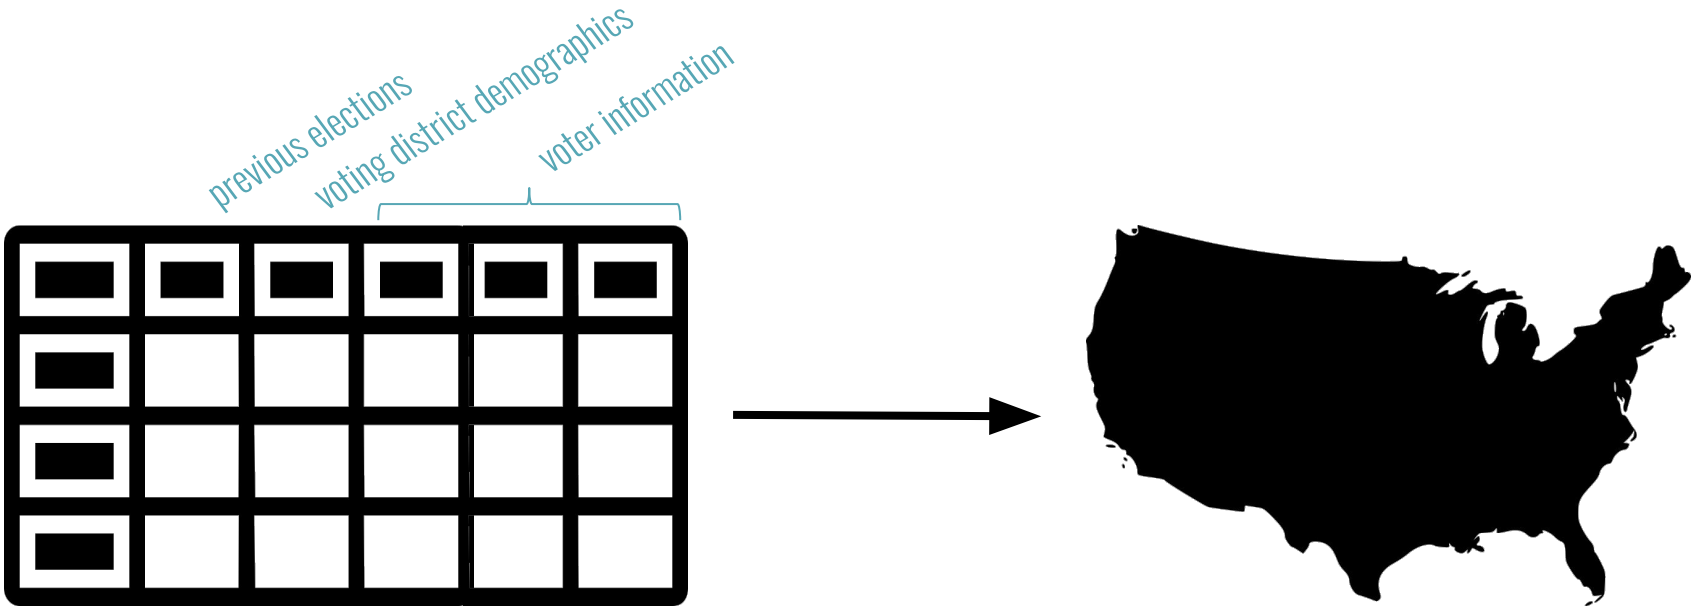

feature selection determines which variables are most predictive and includes them in the model

Note: Occam's Razor

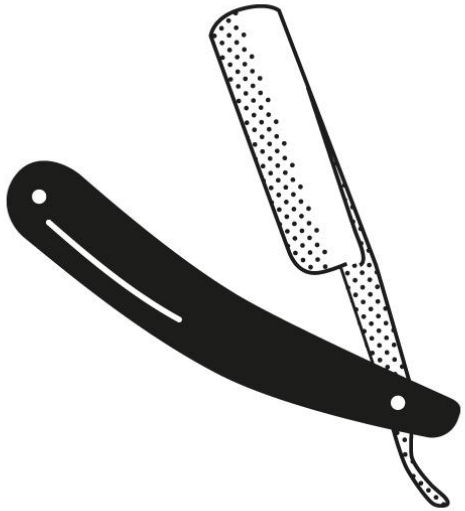

- When faced with two explanations for the *same evidence*, we prefer the explanation that makes the fewest assumptions
- Given *equal performance*, choose the simpler model

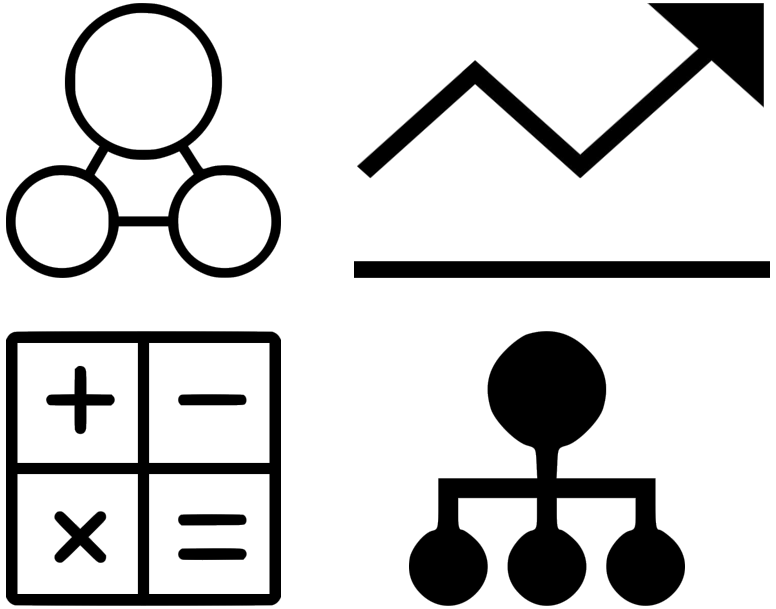

model selection

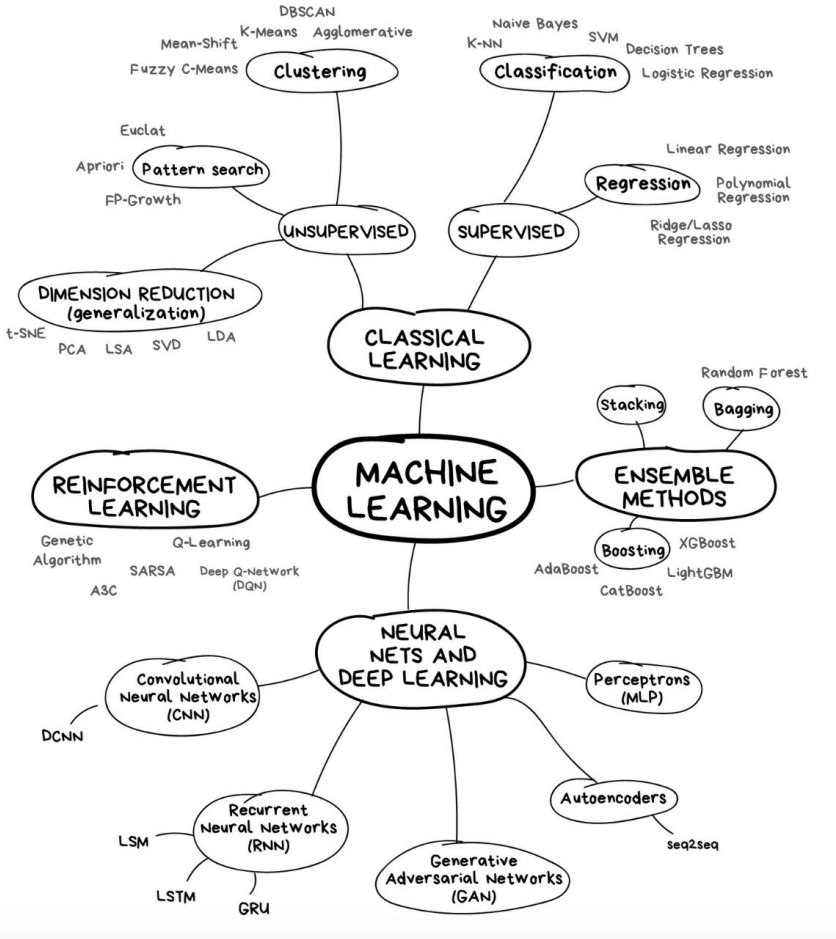

Image source: https://vas3k.com/blog/machine_learning/

Selecting the “Right” Model

- **A** Thought about it and have a thought
- **B** Thought about it and have no thought
- **C** I’m confused

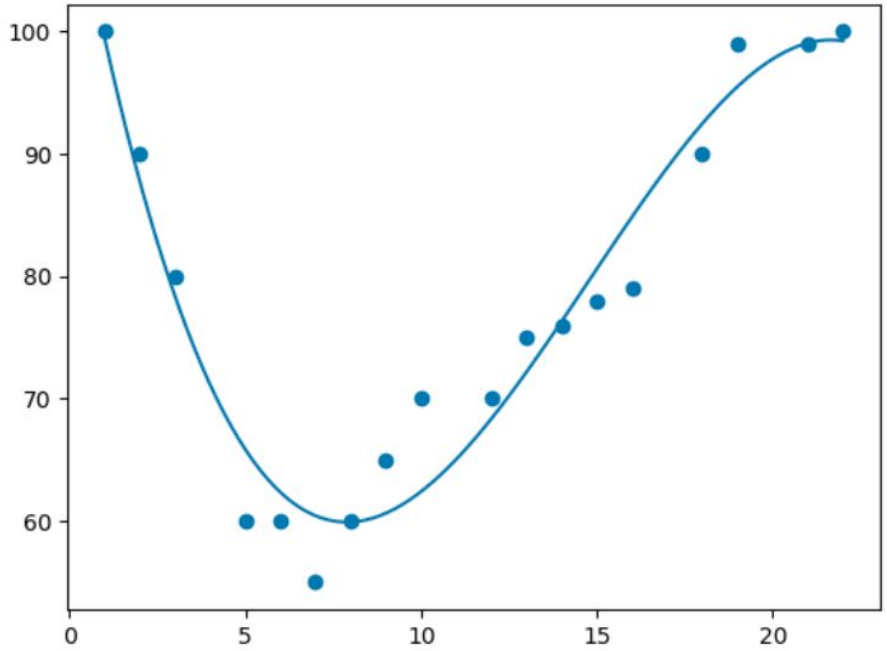

Selecting the “Right” Model

- **A** Thought about it and have a thought
- **B** Thought about it and have no thought
- **C** I’m confused

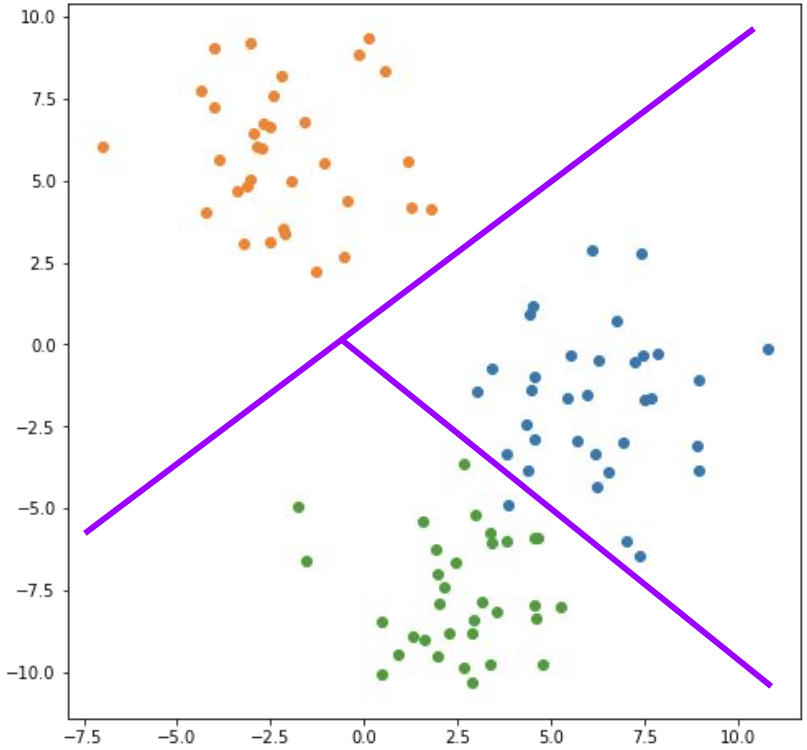

Selecting the “Right” Model

- **A** Thought about it and have a thought
- **B** Thought about it and have no thought
- **C** I’m confused

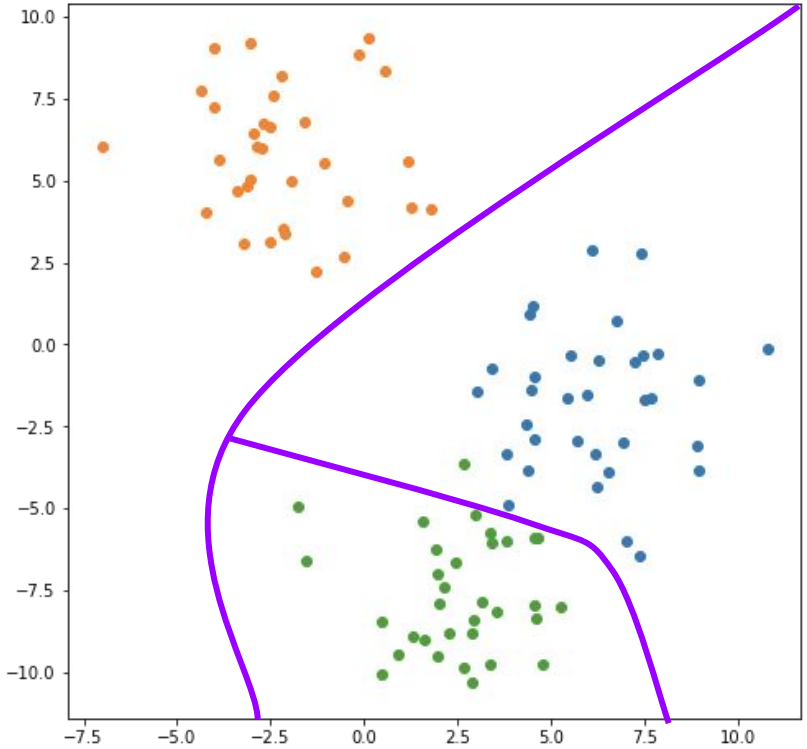

Selecting the “Right” Model

- **A** Thought about it and have a thought
- **B** Thought about it and have no thought
- **C** I’m confused

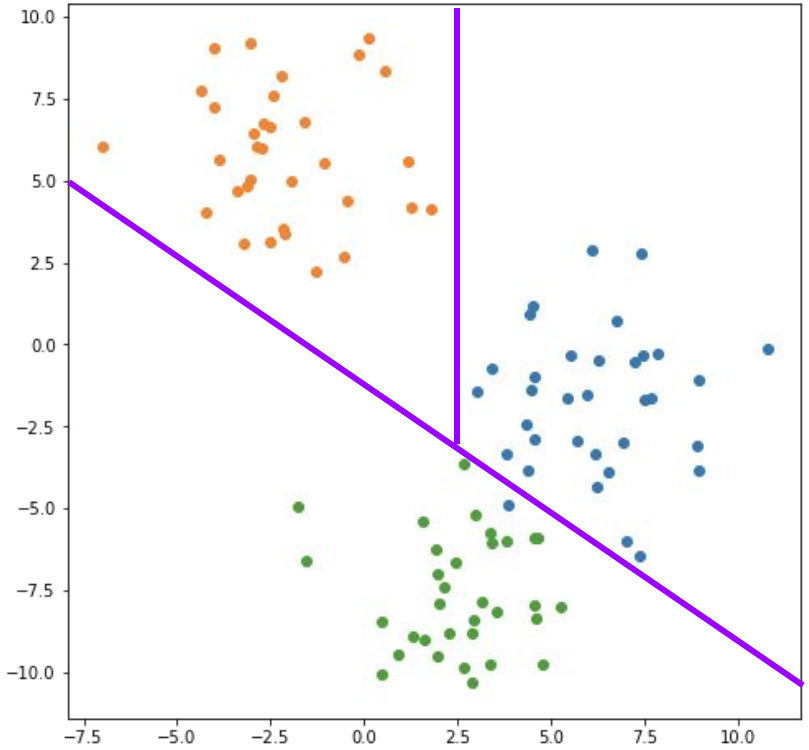

Selecting the “Right” Model

- **A** Thought about it and have a thought
- **B** Thought about it and have no thought
- **C** I’m confused

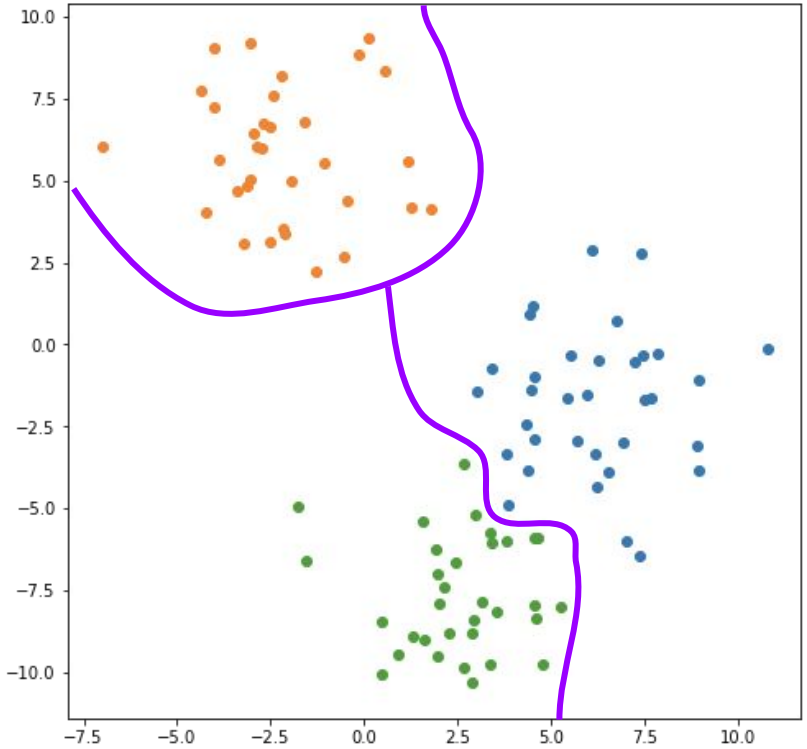

Cross-Validation

In reality, our eyeball meter won't (and sometimes can't) cut it

Real data are messy and live in dimensions we cannot even comprehend (let alone visualize)

Cross-validation offers a systematic way of assessing various models and determine which ones meet our requirements the closest

- Validation Set
- Leave-one-out
- K-Fold

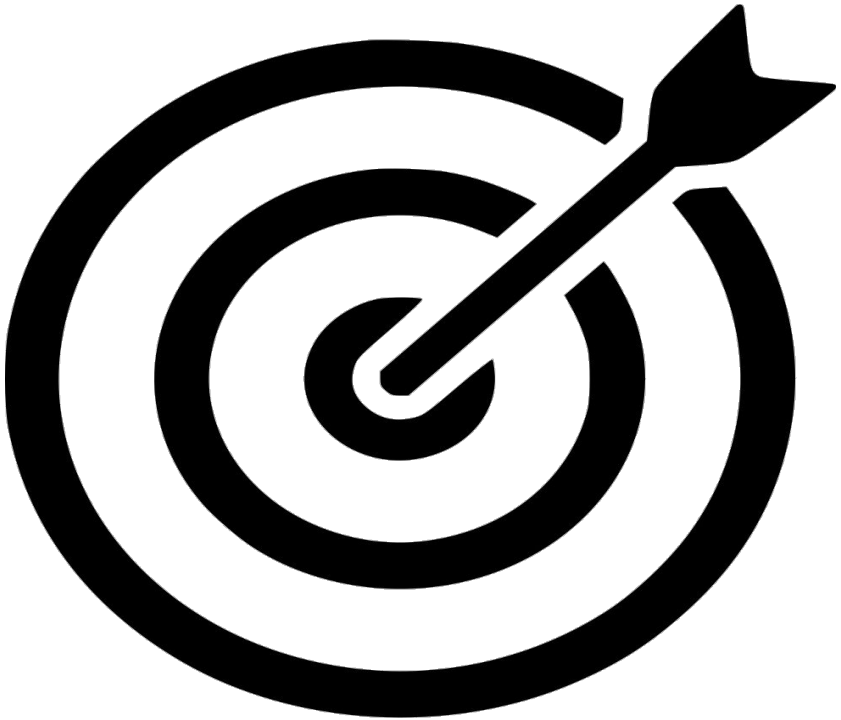

model assessment

Many Metrics!

Assessing machine learning models involves gauging its performance on novel data

- Novel data come from the testing set we left out earlier!
- Performance assessments depend on the specific problem/question
- Different metrics measure different aspects of models
    - e.g. MSE, MAE, accuracy, specificity/sensitivity, etc.

Classical Machine Learning

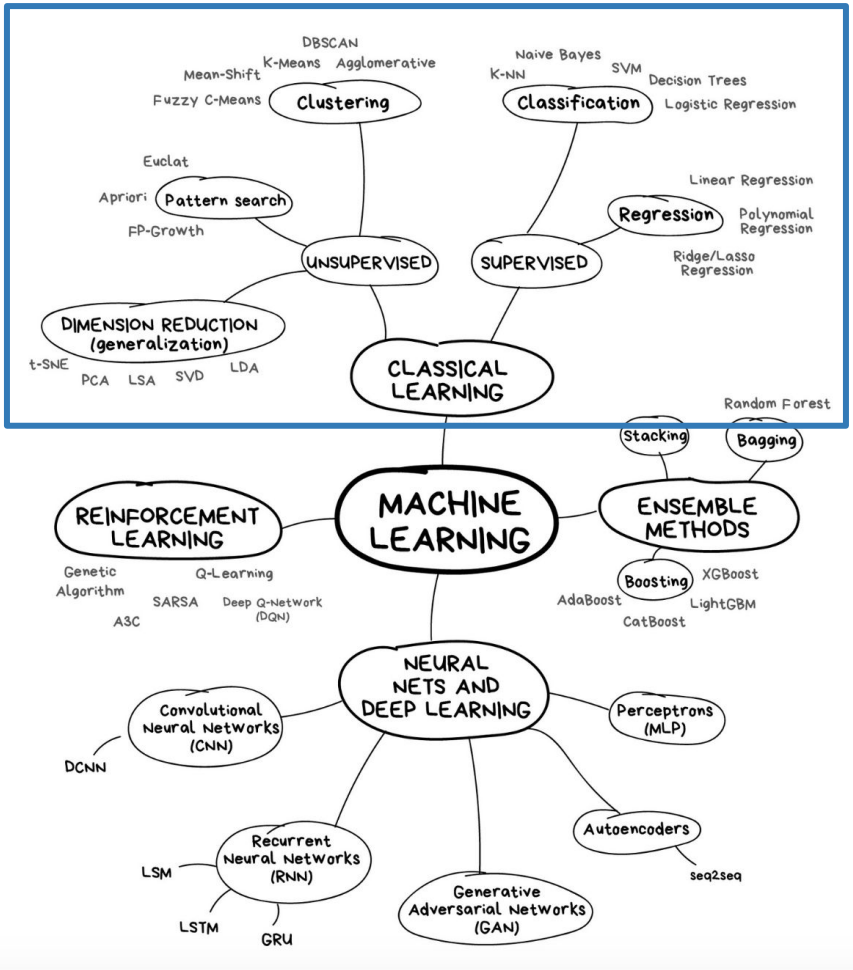

Image source: https://vas3k.com/blog/machine_learning/

CLASSICAL MACHINE LEARNING

Data is pre-categorized or numerical

SUPERVISED

Predict a category

CLASSIFICATION

«Divide the socks by color»

Predict a number

REGRESSION

«Divide the ties by length»

Data is not labeled in any way

UNSUPERVISED

Divide by similarity

CLUSTERING

«Split up similar clothing into stacks»

Find hidden dependencies

DIMENSION REDUCTION (generalization)

«Make the best outfits from the given clothes»

Identify sequences

ASSOCIATION

«Find what clothes I often wear together»

Image source: https://vas3k.com/blog/machine_learning/

Regression

- Predicting a **continuous** outcome by finding a relationship between inputs and outputs
- Makes predictions on new data based on the learned patterns from the training set
- Used for:
  - Stock price forecasts
  - Demand and sales volume analysis
  - Number-time correlations
- Popular regression models include:
  - Linear regression
  - Polynomial regression

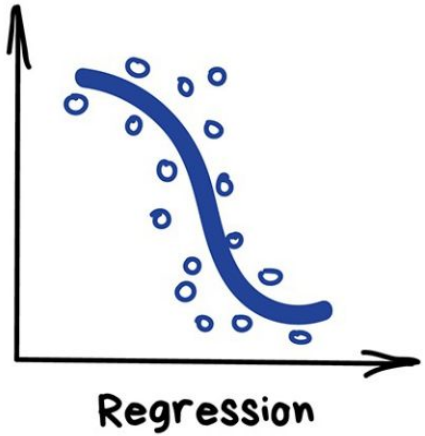

Image source: https://vas3k.com/blog/machine_learning/

PREDICT TRAFFIC JAMS

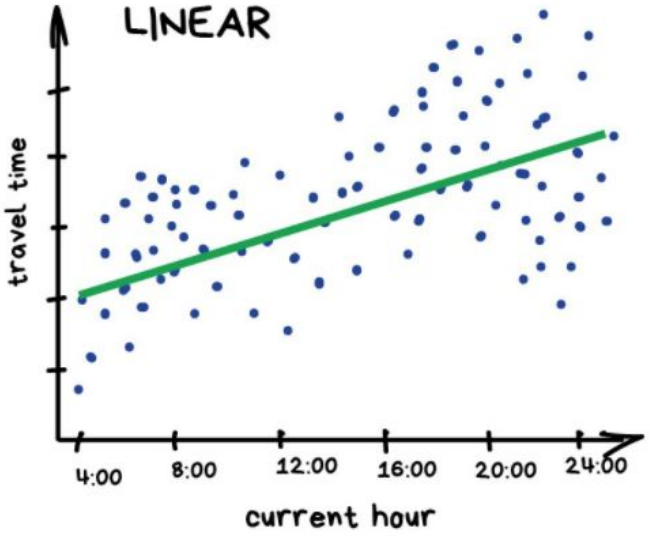

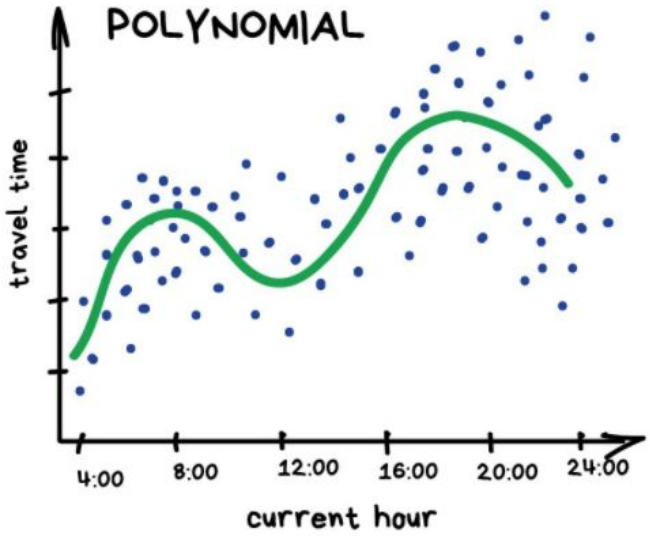

REGRESSION

Image source: https://vas3k.com/blog/machine_learning/

Classification

- Predicting a **categorical** outcome by finding boundaries between classes of data points
- Makes predictions on new data based on where within the decision boundary it falls
- Used for:
  - Spam filtering
  - Fraud detection
  - Language detection
- Popular classification models include:
  - Logistic regression
  - Support Vector Machine
  - Decision Tree

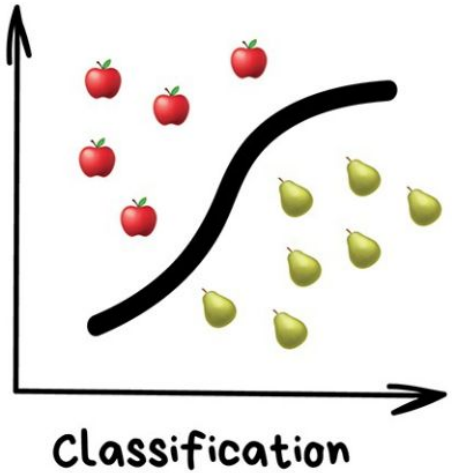

Image source: https://vas3k.com/blog/machine_learning/

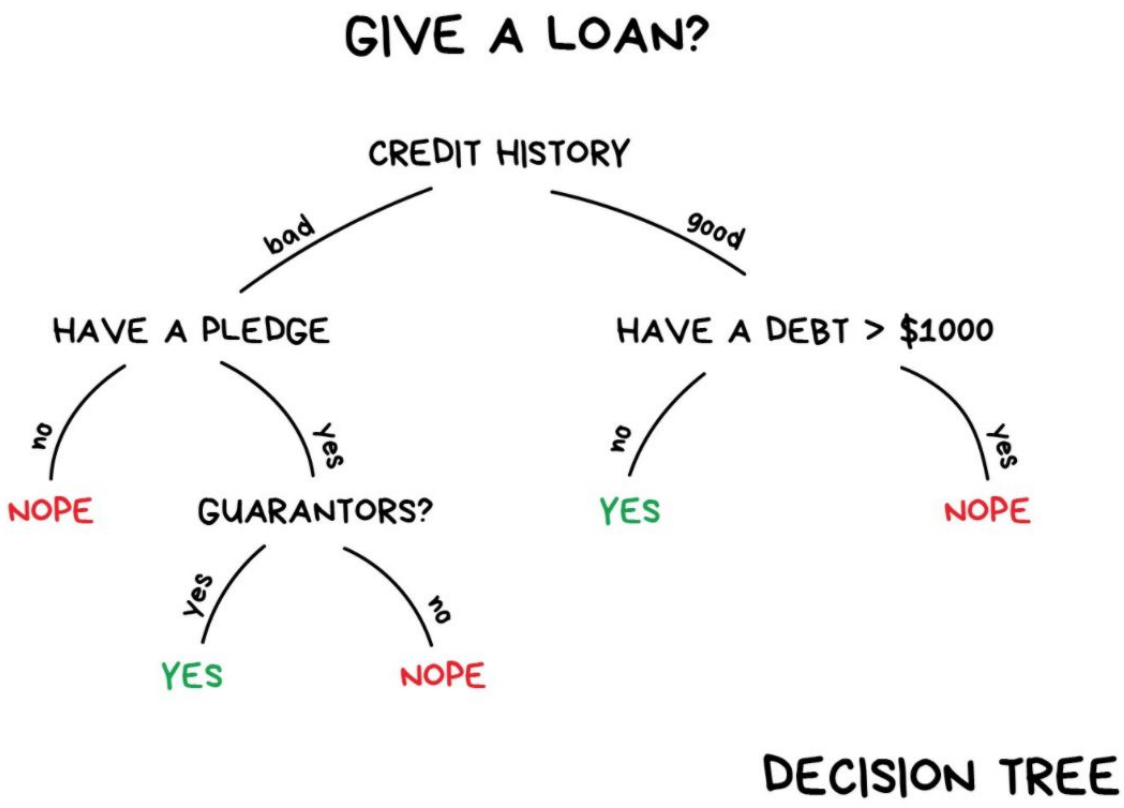

Image source: https://vas3k.com/blog/machine_learning/

SEPARATE TYPES OF ALCOHOL

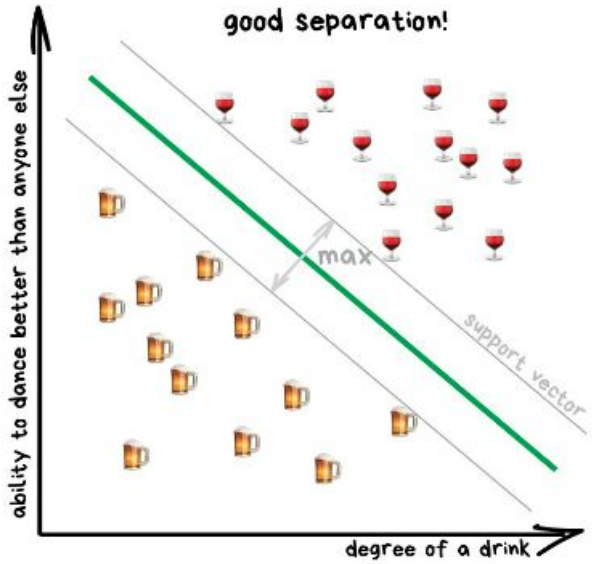

SUPPORT VECTOR MACHINE

Image source: https://vas3k.com/blog/machine_learning/

Supervised vs. Unsupervised

Supervised algorithms use labeled data → Checks answers and improves over time

- Learning relationships between inputs and outputs to make predictions
- Goal is to minimize error or maximize accuracy in predictions

Unsupervised algorithms use unlabeled data → No “correct” answers

- Commonly used to discover new patterns where they are unknown → research!!
- Verifies itself with similarity/stability scores

Regression Walkthrough

Can we determine what that function (`f`) *is* using these data?

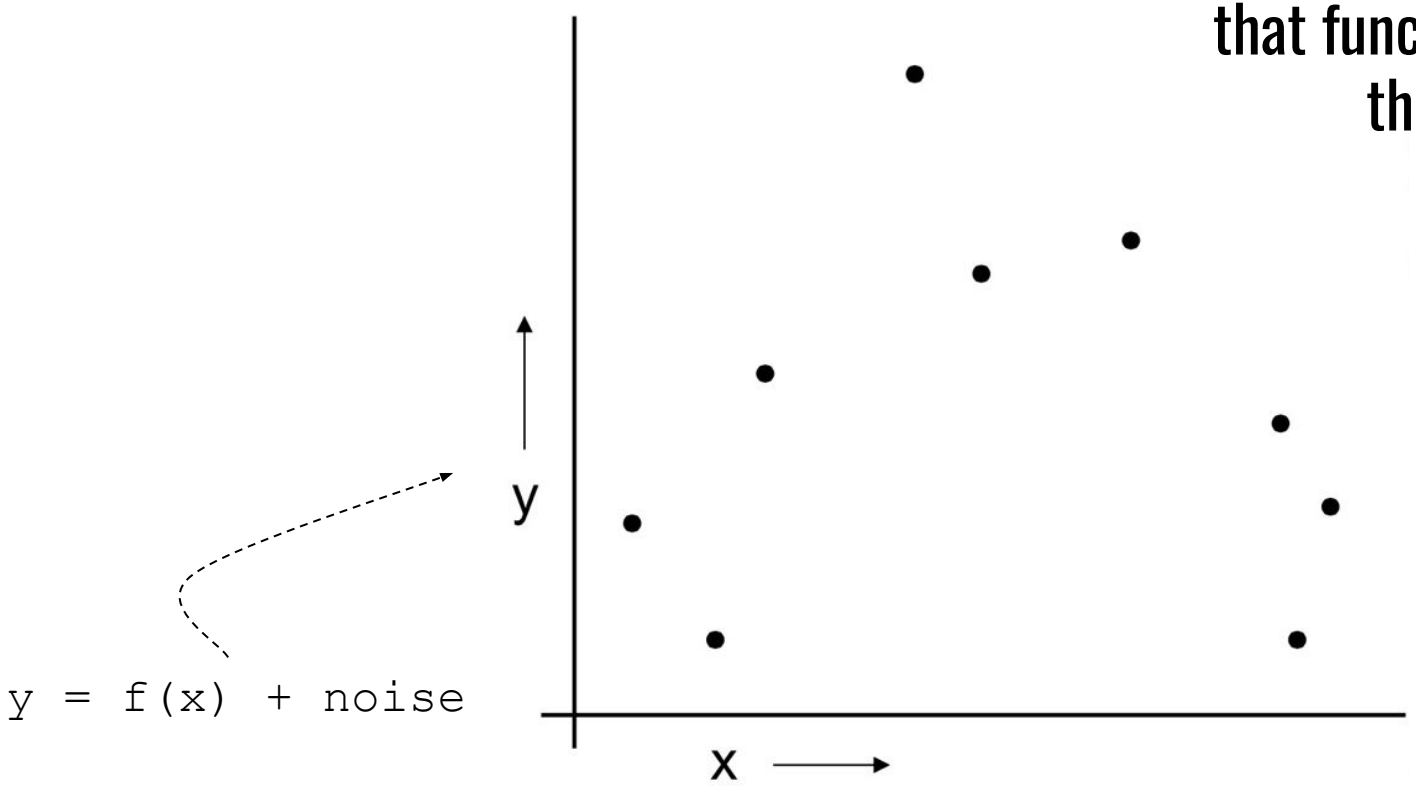

Linear regression

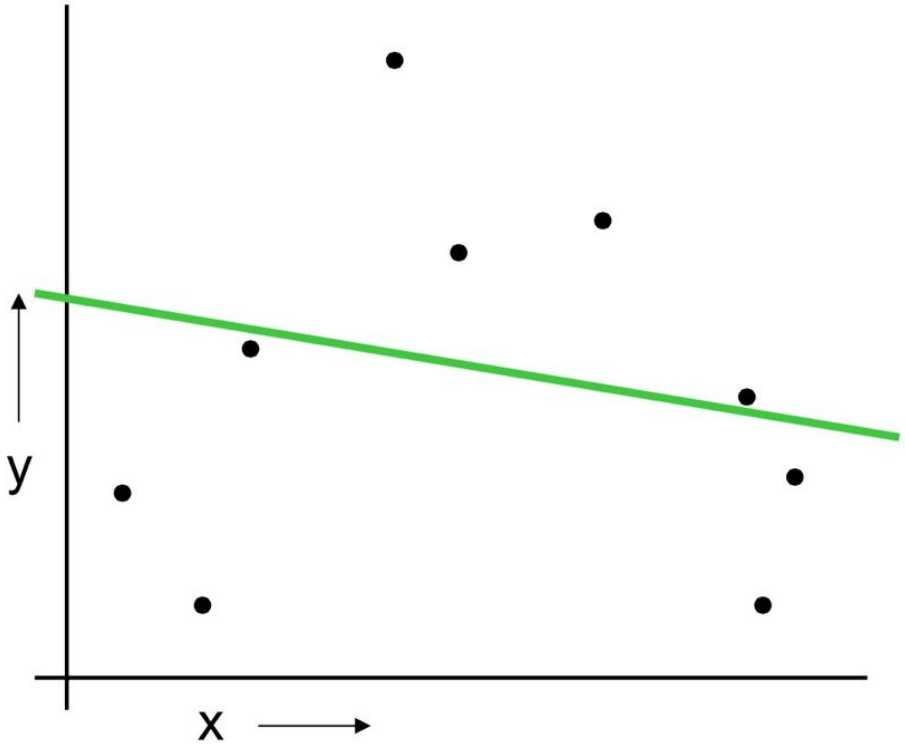

Quadratic regression

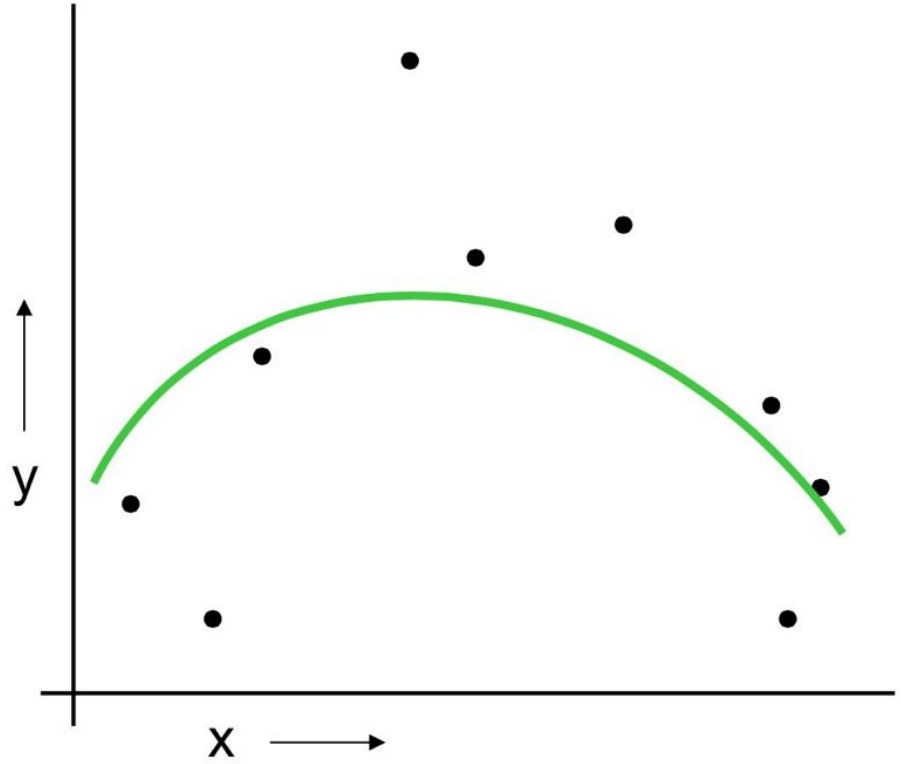

Piecewise linear nonparametric regression

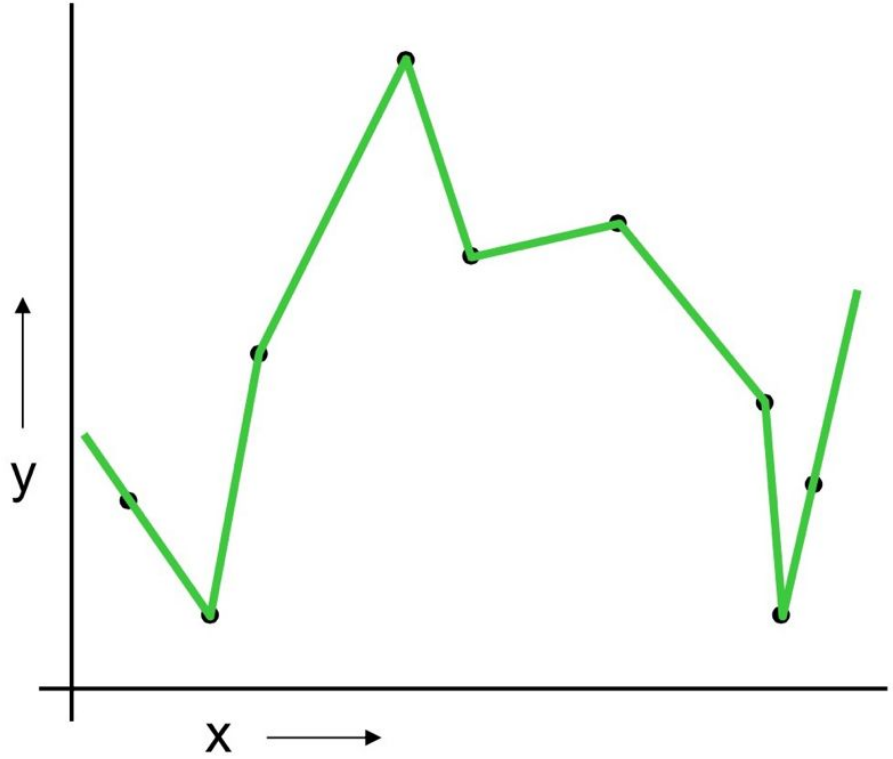

Which to choose?

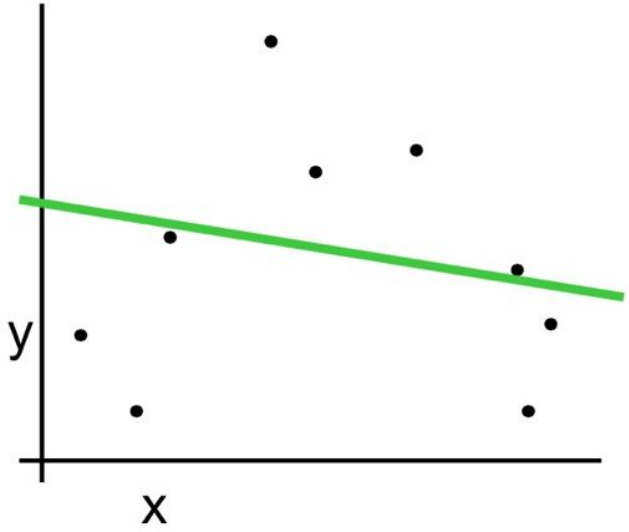

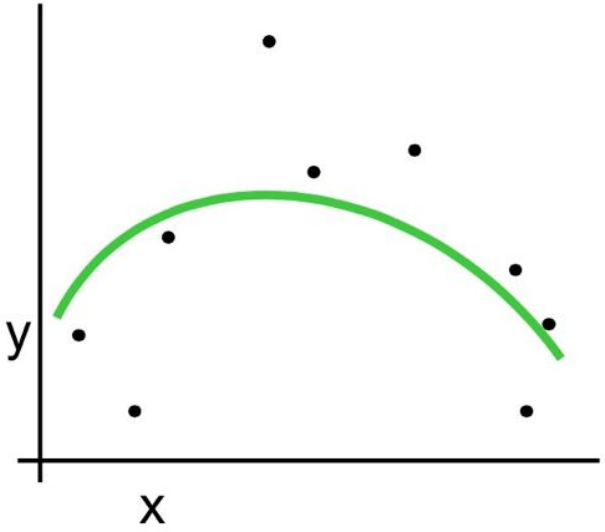

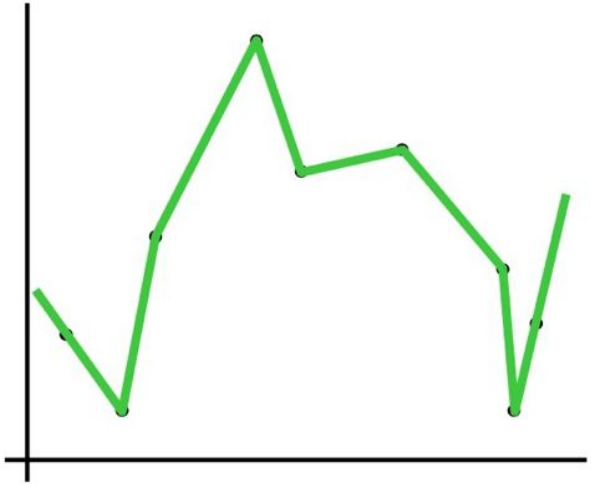

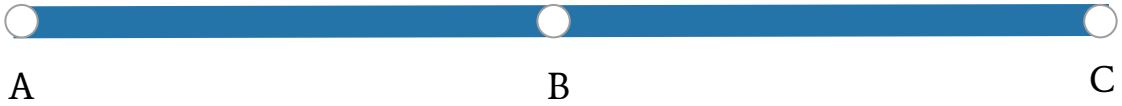

The data partition method

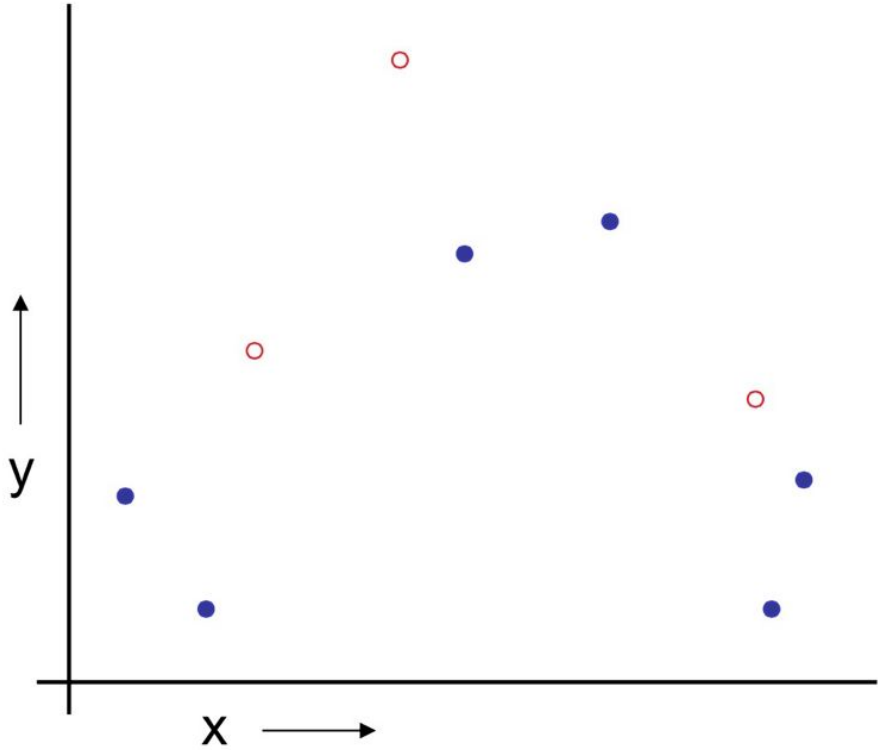

1. Randomly choose 30% of the data to be in a test set

2. The remainder is a training set

Train the model on your training set

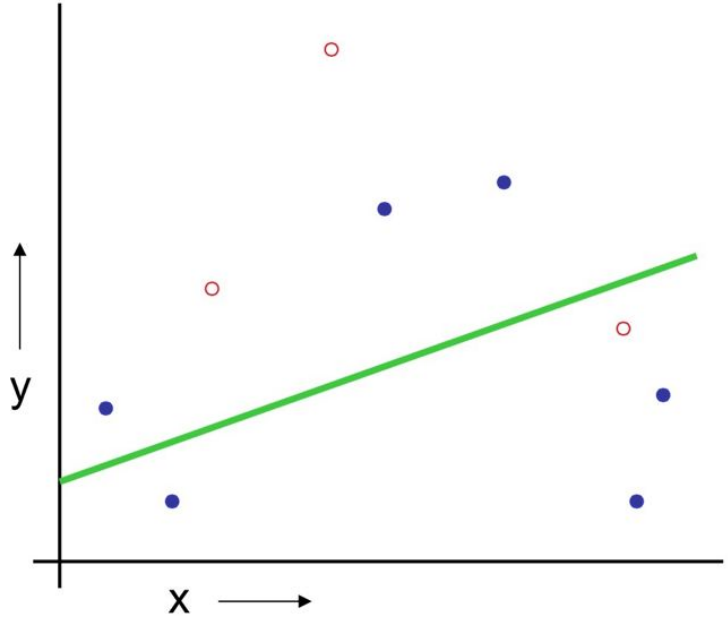

(Linear regression example)

1. Randomly choose 30% of the data to be in a test set

2. The remainder is a training set

3. Perform your regression on the training set

Assess future performance using the test set

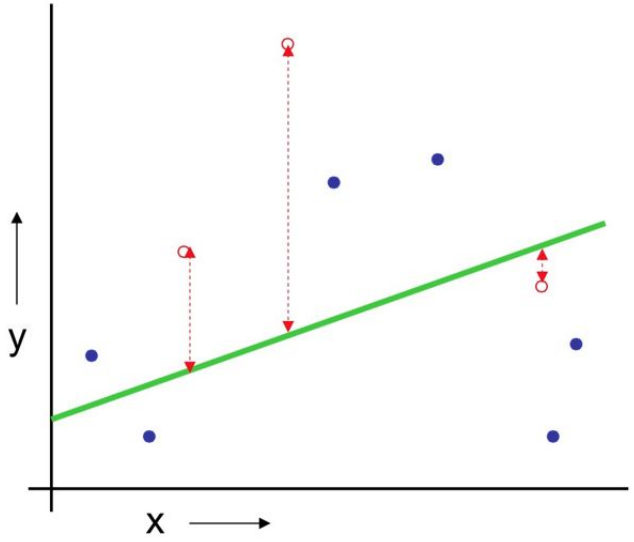

(Linear regression example)

Mean Squared Error = 2.4

1. Randomly choose 30% of the data to be in a test set

2. The remainder is a training set

3. Perform your regression on the training set

4. Estimate your future performance with the test set

Go through this process for each possible model

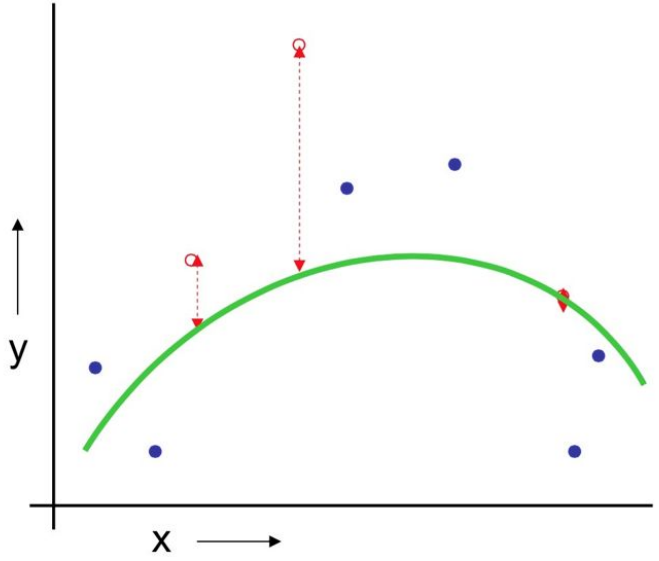

(Quadratic regression example)

Mean Squared Error = 0.9

1. Randomly choose 30% of the data to be in a test set

2. The remainder is a training set

3. Perform your regression on the training set

4. Estimate your future performance with the test set

Go through this process for each possible model

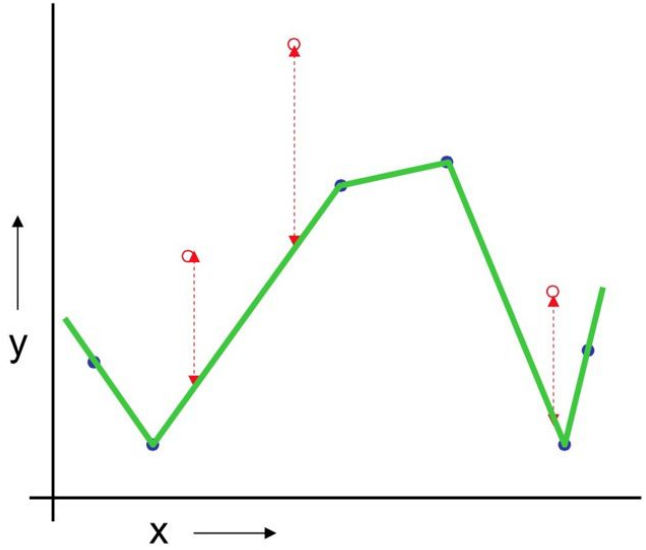

(Join the dots example)

Mean Squared Error = 2.2

1. Randomly choose 30% of the data to be in a test set

2. The remainder is a training set

3. Perform your regression on the training set

4. Estimate your future performance with the test set

Pros and cons of data partitioning

Pros:

- Simple approach
- Can choose model with best test-set score

Cons:

- Model fit on 30% less data than you have
- Without a large data set, removing 30% of the data could bias prediction

Leave one out cross validation (LOOCV)

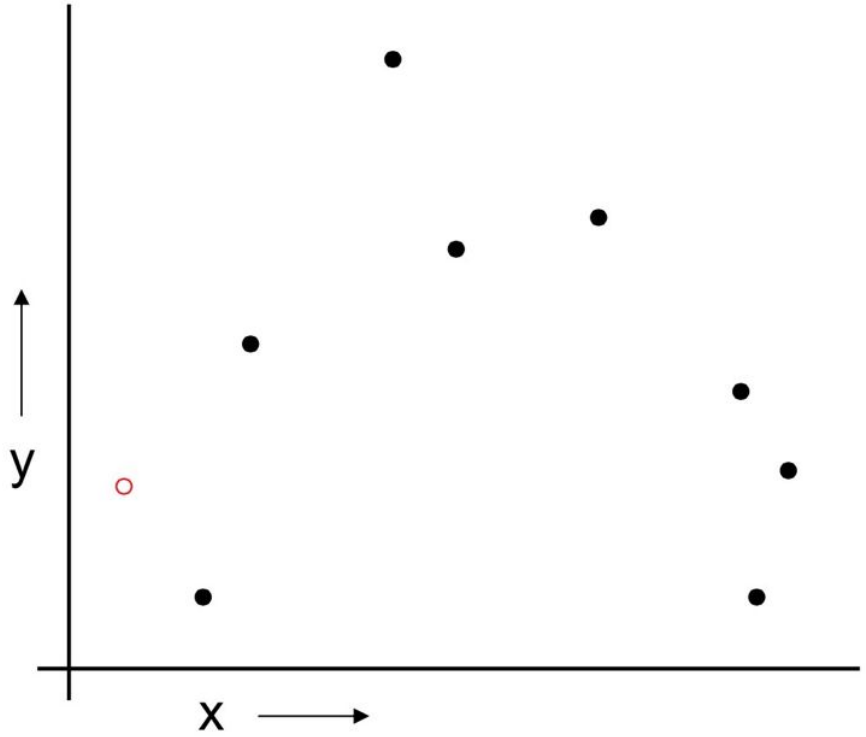

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record

Leave one out cross validation (LOOCV)

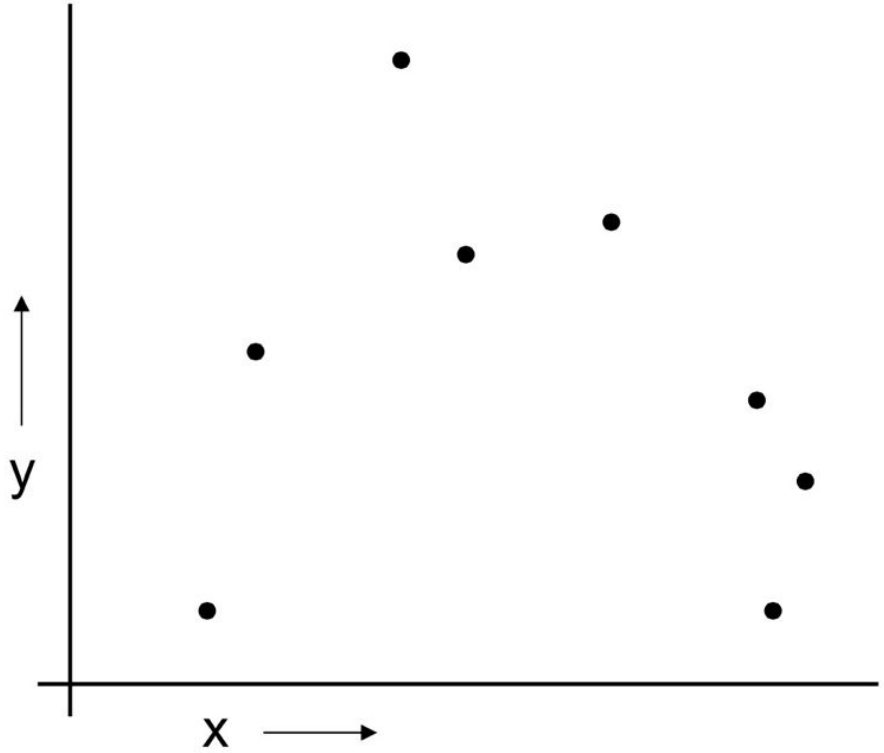

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset

Leave one out cross validation (LOOCV)

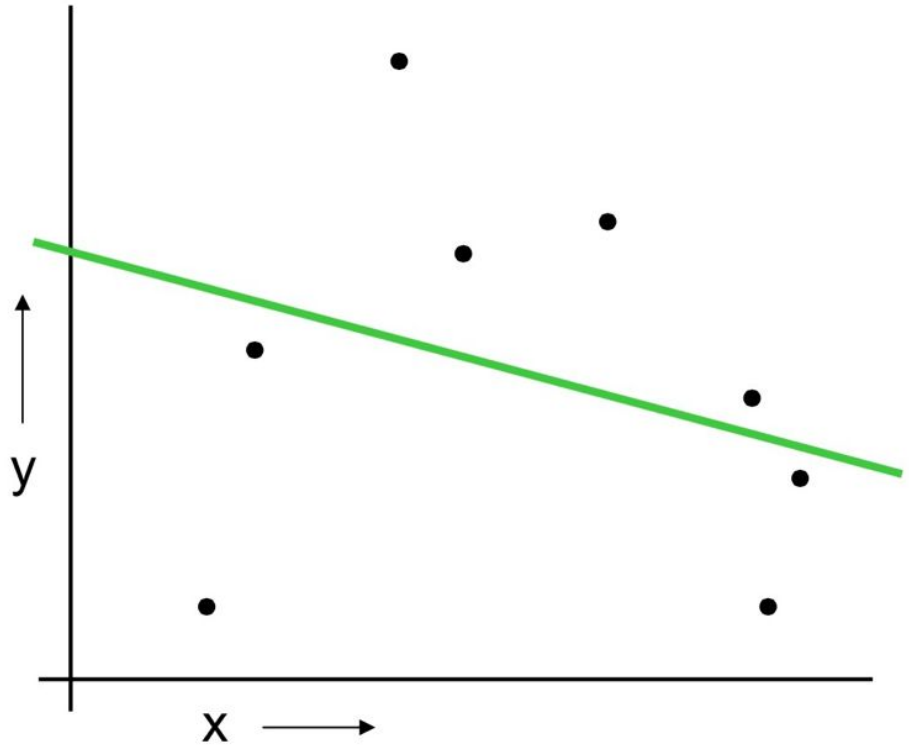

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining R-1 datapoints

Leave one out cross validation (LOOCV)

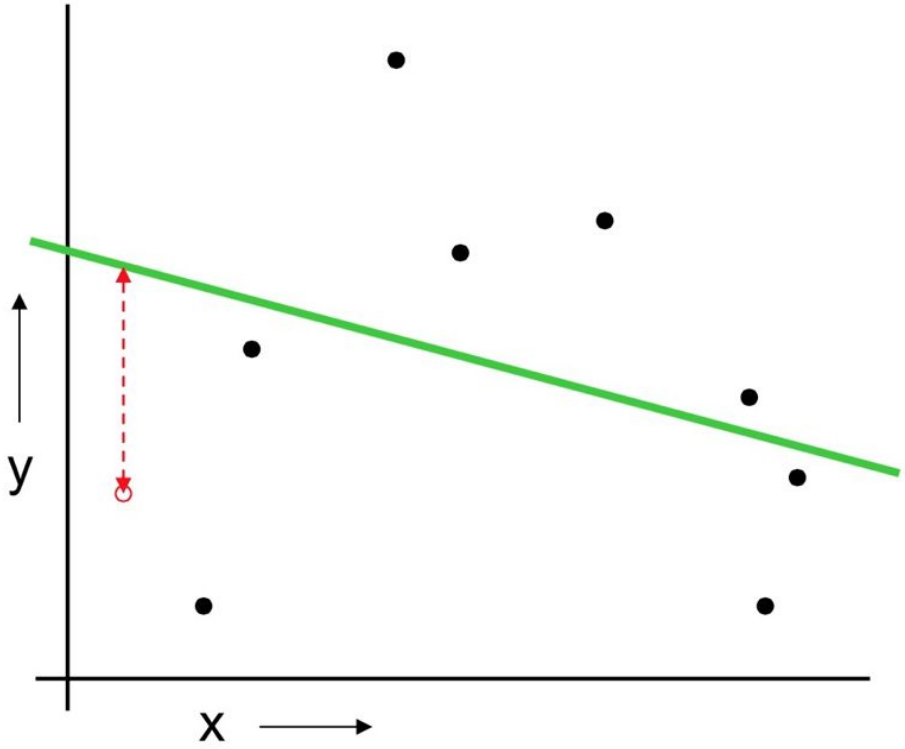

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining R-1 datapoints
4. Note your error (x_k, y_k)

Leave one out cross validation (LOOCV)

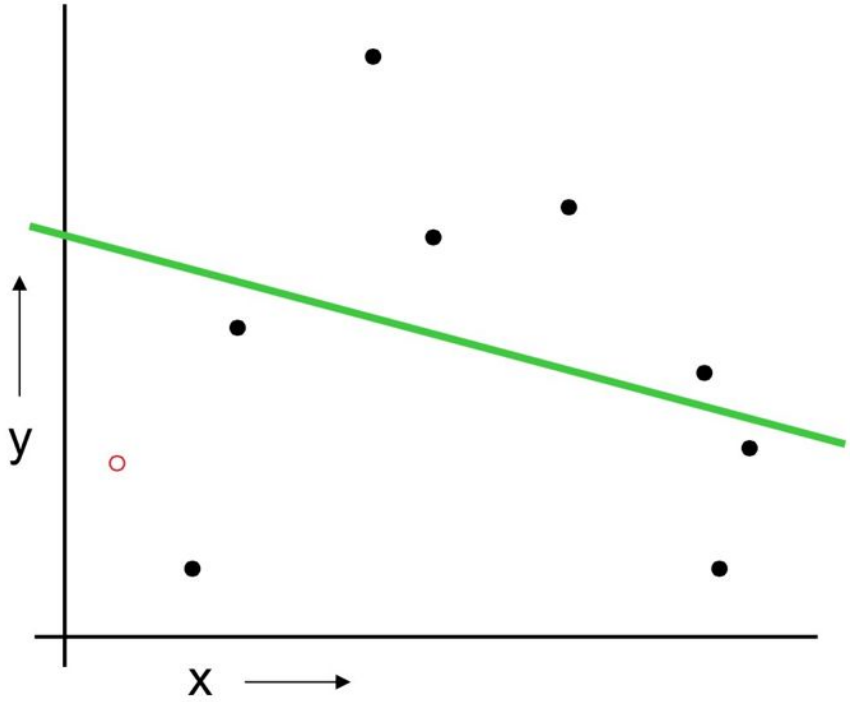

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining R-1 datapoints
4. Note your error (x_k, y_k)

When you've done all points, report the mean error.

Leave one out cross validation (LOOCV)

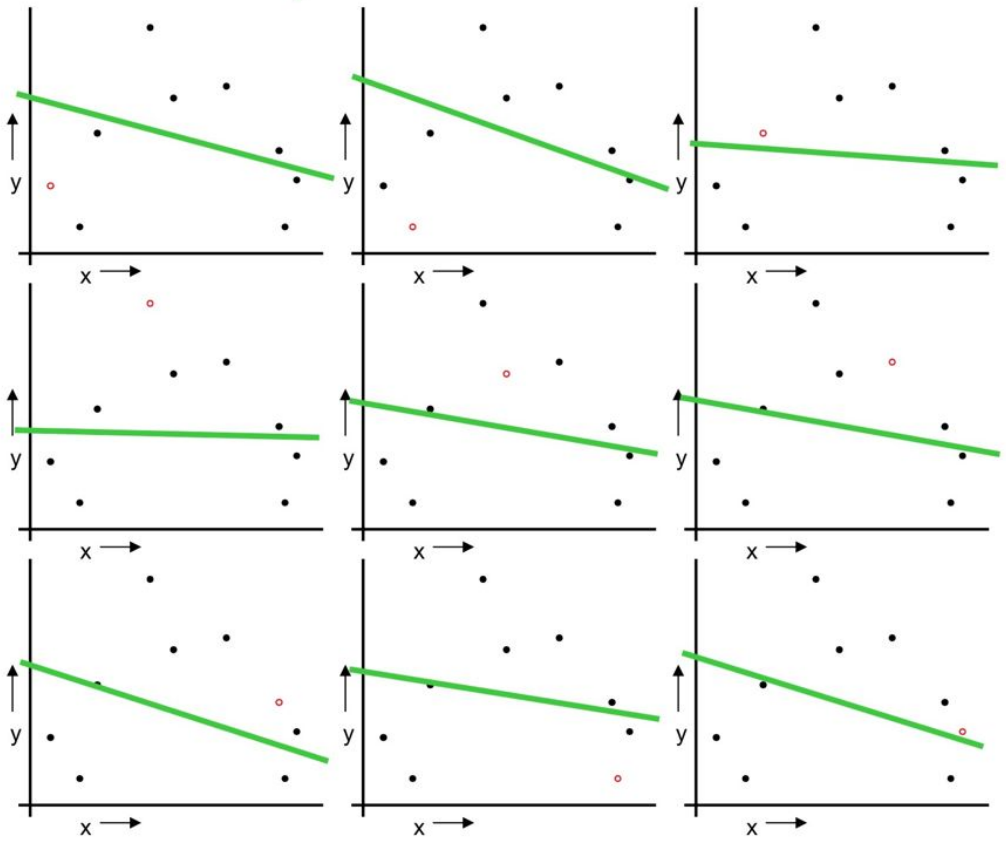

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining R-1 datapoints
4. Note your error (x_k, y_k)

When you've done all points, report the mean error.

$MSE_{LOOCV} = 2.12$

Leave one out cross validation (LOOCV)

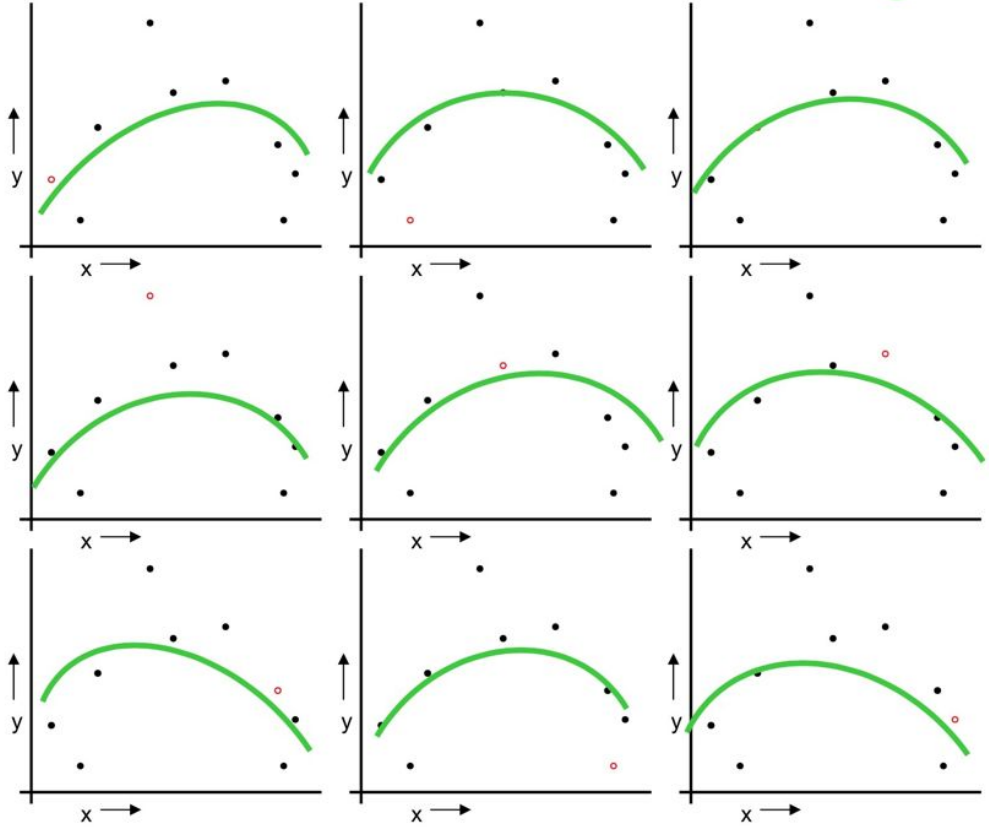

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining R-1 datapoints
4. Note your error (x_k, y_k)

When you've done all points, report the mean error.

$MSE_{LOOCV} = 0.962$

Leave one out cross validation (LOOCV)

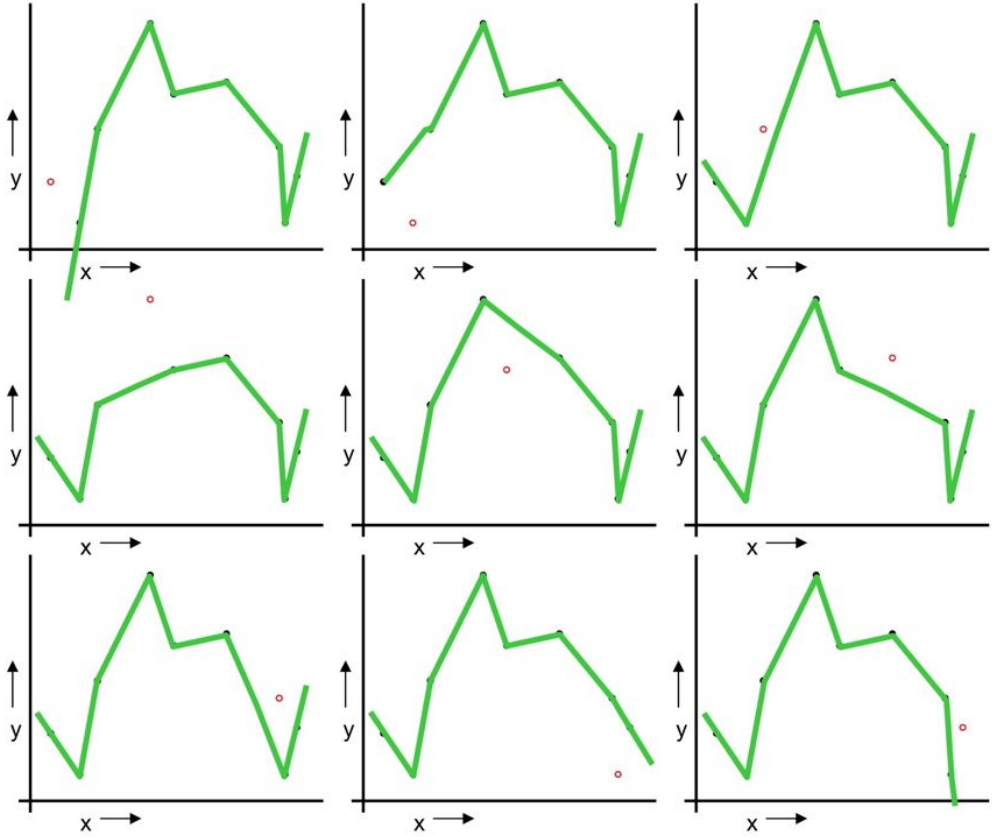

For k=1 to R

1. Let (x_k, y_k) be the k^{th} record
2. Temporarily remove (x_k, y_k) from the dataset
3. Train on the remaining R-1 datapoints
4. Note your error (x_k, y_k)

When you've done all points, report the mean error.

$MSE_{LOOCV} = 3.33$

Method Comparison

	Cons	Pros
Data partitioning	Variance: unreliable estimate of future performance	Cheap
LOOCV	Computationally expensive; has weird behavior	Uses all your data

k-fold cross validation

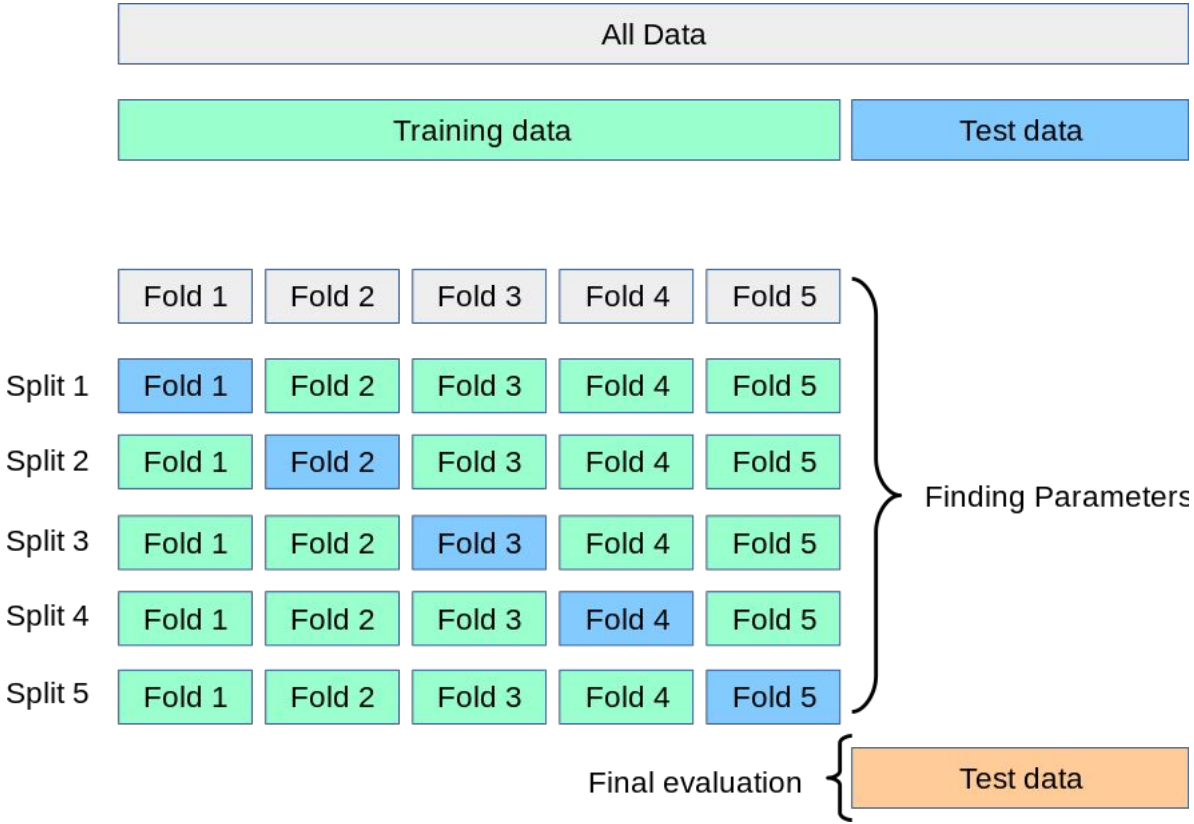

https://scikit-learn.org/stable/modules/cross_validation.html

k-fold cross validation

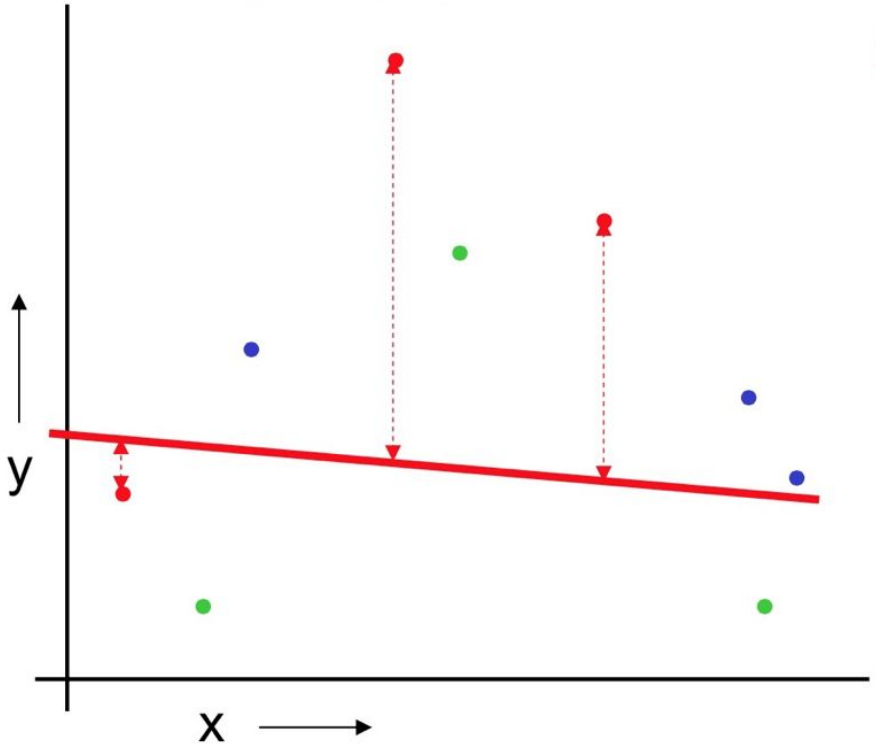

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

k-fold cross validation

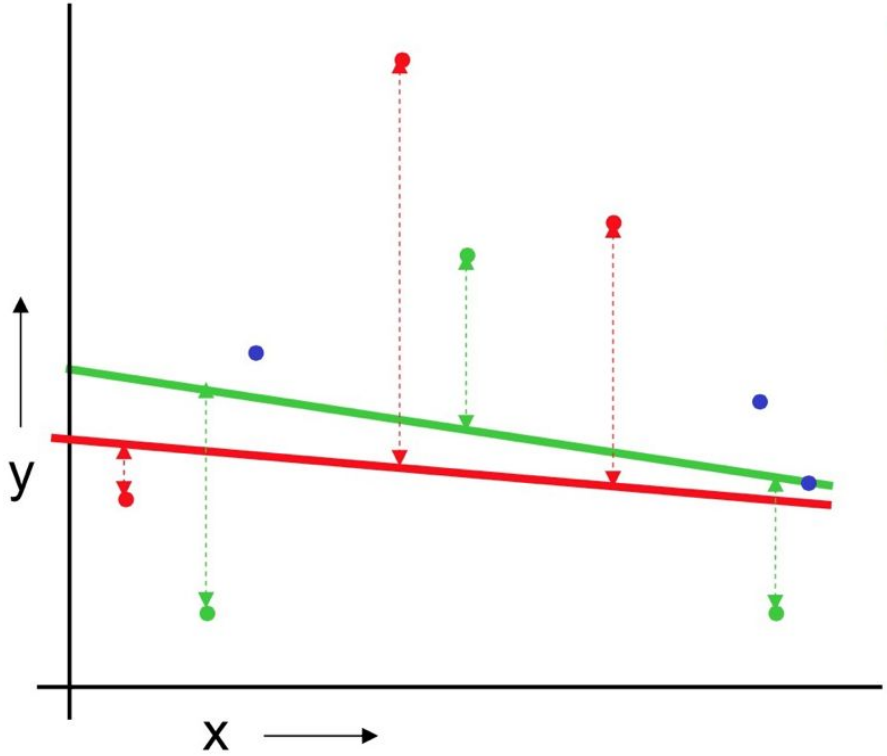

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

k-fold cross validation

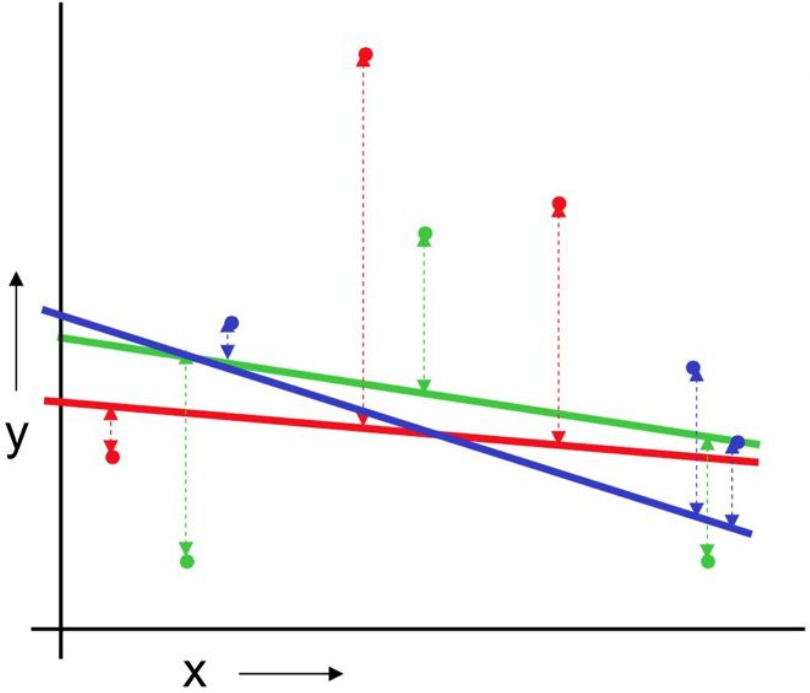

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

k-fold cross validation

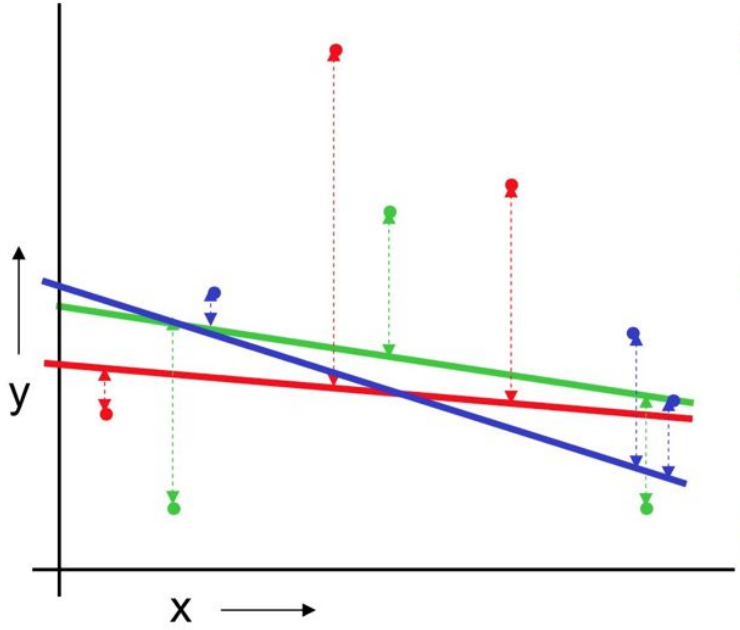

Linear Regression

$MSE_{3FOLD}=2.05$

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

Then report the mean error

k-fold cross validation

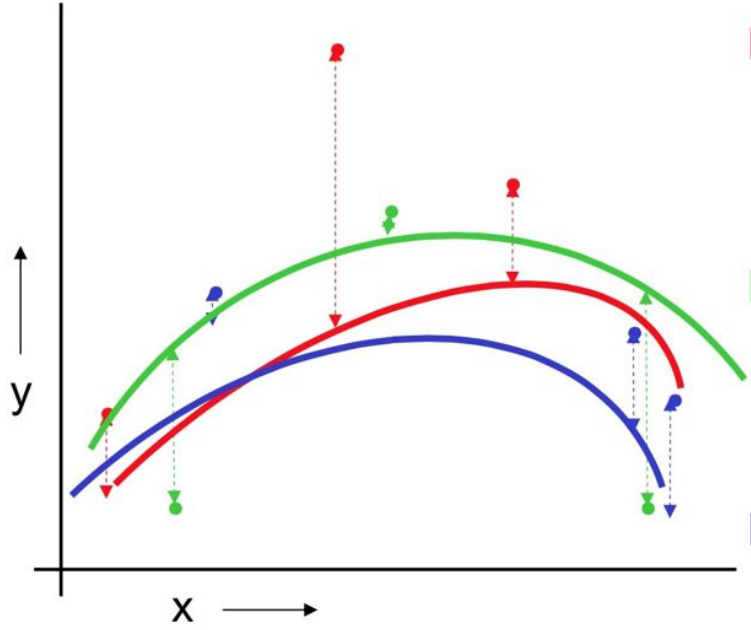

Quadratic Regression

$MSE_{3FOLD}=1.11$

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

Then report the mean error

k-fold cross validation

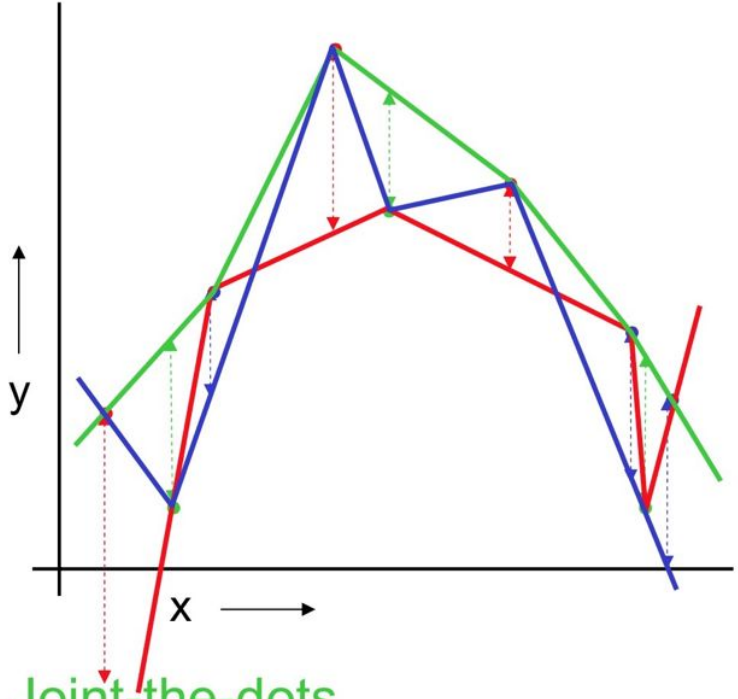

Joint-the-dots

$MSE_{3FOLD}=2.93$

For the red partition: Train on all the points not in the red partition. Find the test-set sum of errors on the red points.

For the green partition: Train on all the points not in the green partition. Find the test-set sum of errors on the green points.

For the blue partition: Train on all the points not in the blue partition. Find the test-set sum of errors on the blue points.

Then report the mean error

k-fold cross validation

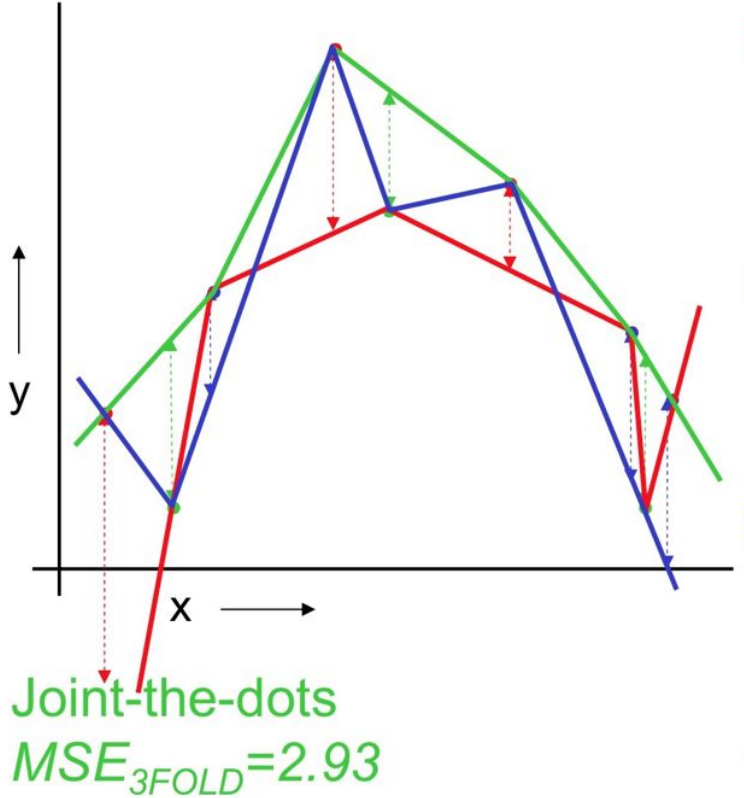

NOTE: Notice each fold generates a dramatically different model than the others?

- This model has *overfit* the training data
- i.e. You only memorized the study guide
- Perfect training performance, but poor predictive power on test set

Validator

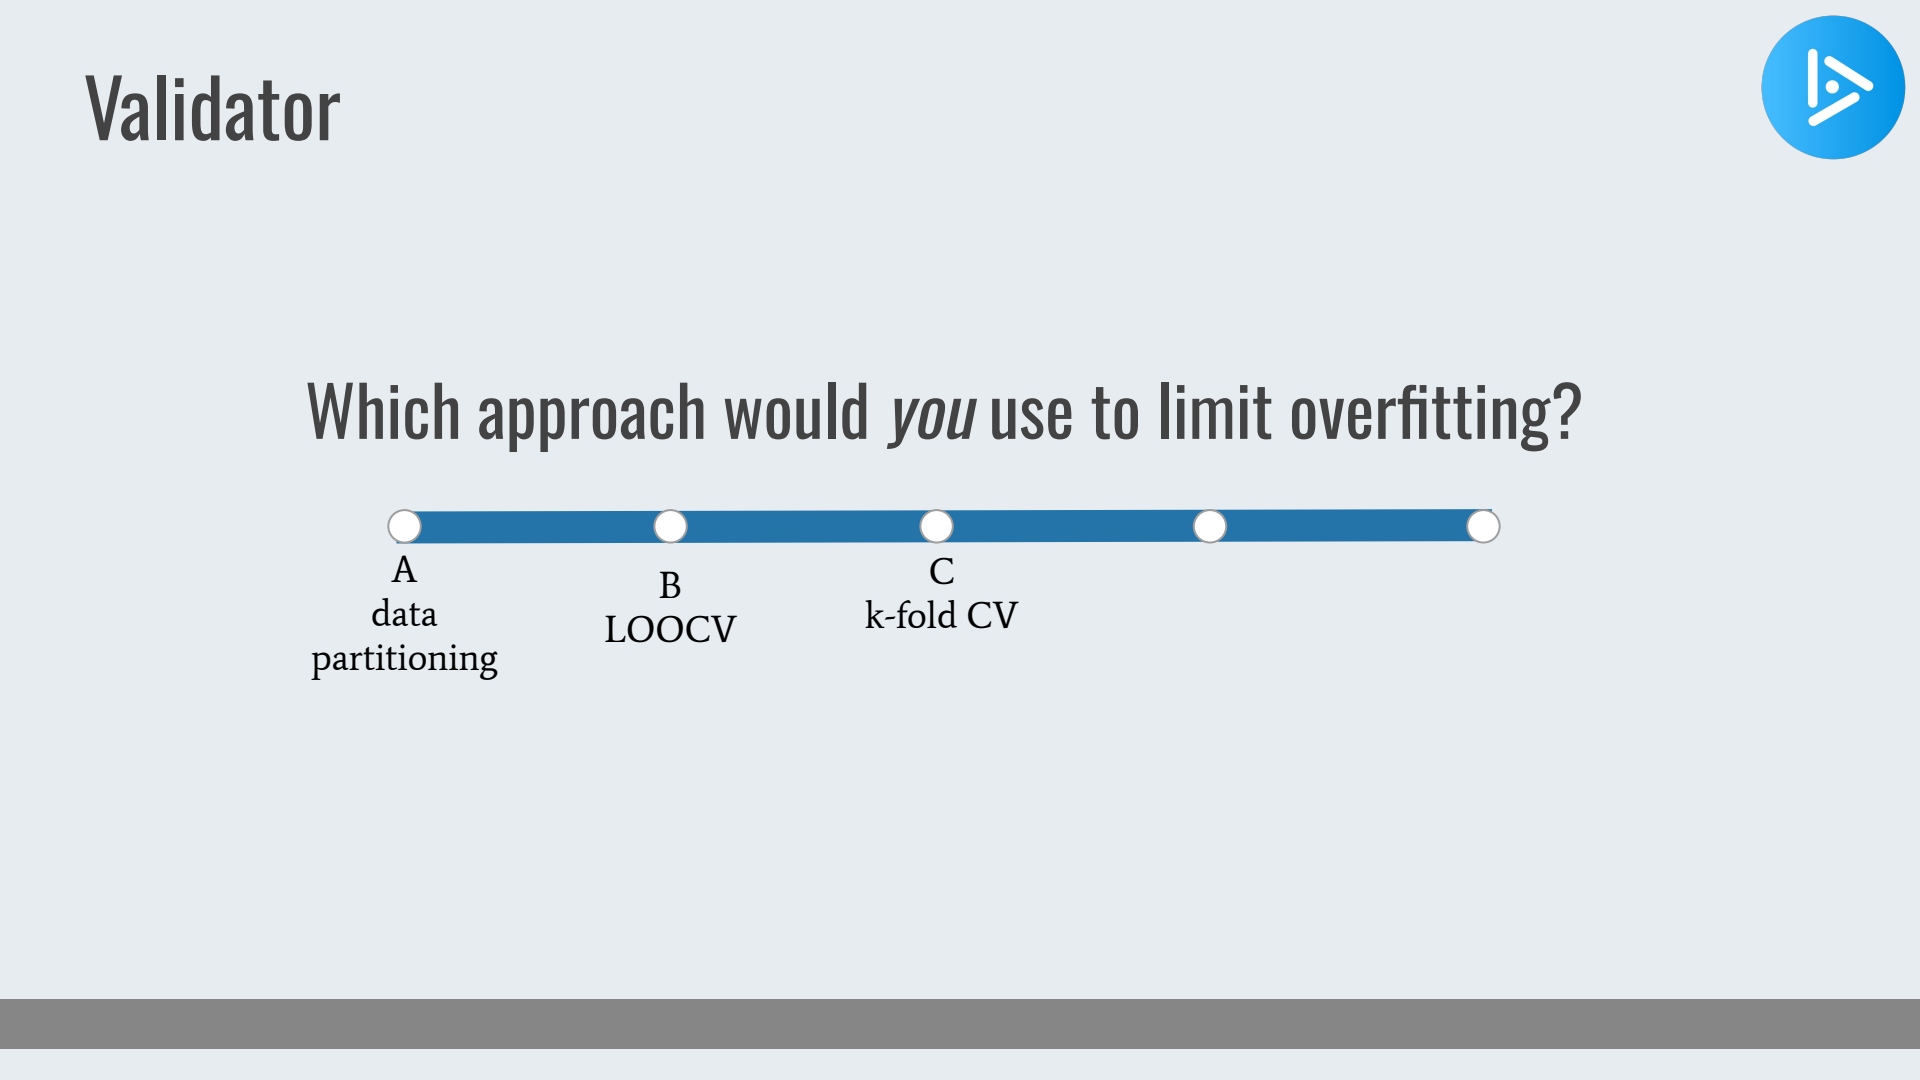
Which approach would *you* use to limit overfitting?
A
data
partitioning
B
LOOCV
C
k-fold CV

Validator

Given the example we just worked, how would *you* model these data?