

How to Be Wrong

Shannon E. Ellis, Ph.D
UC San Diego



Department of Cognitive Science
sellis@ucsd.edu

How Charts Lie



Getting Smarter about
Visual Information

Alberto Cairo

2019


Now you, too, can out-double-talk your accountant;
confuse your political opponent;

How to
prove that your product has secret built-in goodness!

LIE
with
Statistics

by Darrell Huff

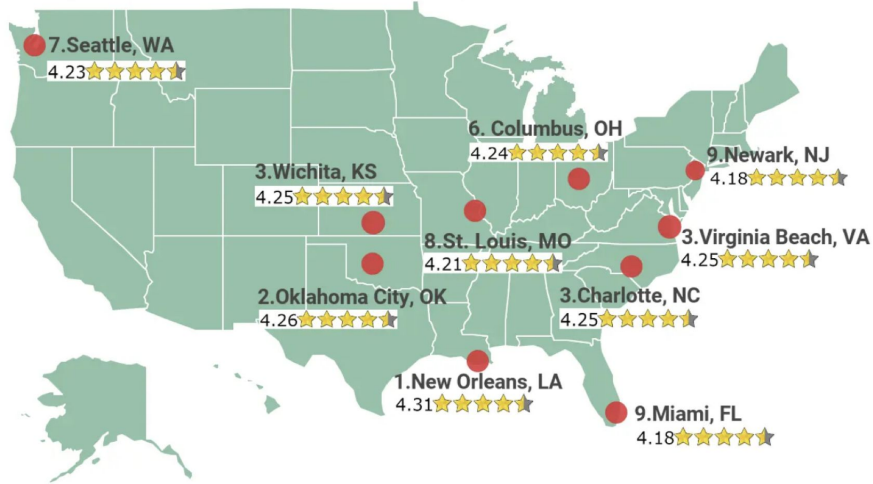
PICTURES BY
Irving Geis



1954

Top Cities For BBQ in the U.S.

An analysis of TripAdvisor restaurant reviews by chefspencil.com



Worst Cities For BBQ in the U.S.


An analysis of TripAdvisor restaurant reviews by chefspencil.com




- ☐ A Thought about it and have a thought
- ☐ B Thought about it and have no thought
- ☐ C I'm confused

Errors of measurement


**Accurate
and precise**



**Accurate
but not precise**



**Precise
but not accurate**




Proxy Measurements

For when we can't measure directly what we care about, but we can measure something related

- Unemployment rate → our general economy
- Gross domestic product → standard of living
- BMI → health
- Your suggestions?
- Grades → learning


**Do we understand what
we are measuring?**

Survivorship bias




Abraham Wald

Survivorship Bias



Where would you put the armor?



Are we measuring what's relevant?

Academic Life at University of California--San Diego

The student-faculty ratio at University of California--San Diego is 19:1, and the school has 47.4% of its classes with fewer than 20 students. The most popular majors at University of California--San Diego include: Biology, General; Mathematics; Economics; International/Global Studies; and Computer Science. The average freshman retention rate, an indicator of student satisfaction, is 94%.

Class Sizes



Classes with fewer than 20 students 47.4% 20-49 25.8% 50 or more 26.8%

| | |
|-----------------------|------|
| Student-faculty ratio | 19:1 |
|-----------------------|------|

| | |
|------------------------|-----|
| 4-year graduation rate | 65% |
|------------------------|-----|

- ☒ **A** This reflects my experience
- ☐ **B** This kind of reflects my experience
- ☐ **C** This does not reflect my experience


| | | % of classes with this many students | Cumulative % | Fraction of classes with this many students * min number of students in that class type | % of students in these classes (normalized version of column to the left) | |
|--|---|--------------------------------------|--------------|---|---|--|
| | 2-9 students: | 12% | 12% | 0.24 | 0.67% | |
| | 10-19 students | 32% | 44% | 3.2 | 8.95% | |
| Median class size as experienced by faculty | 20-29 students: | 14% | 58% | 2.8 | 7.83% | |
| | 30-39 students: | 8% | 66% | 2.4 | 6.72% | |
| | 40-49 students: | 4% | 70% | 1.6 | 4.48% | |
| | 50-99 students: | 11% | 81% | 5.5 | 15.39% | |
| Median class size as experienced by students | Over 100 students: | 20% | 101% | 20 | 55.96% | |
| | | | Sum: | 35.74 | | |
| | | | | | | |
| | Data from https://www.collegedata.com/college/University-of-California-San-Diego/ | | | | | |

The Lamppost Problem

aka. Streetlight Effect

aka. The Drunkard's Search

aka. Picking a crap proxy



Goodhart's law

Once a metric becomes a goal, you're hosed

Volkswagen emissions scandal

🌐 26 languages ▾

Article [Talk](#) [Read](#) [Edit](#) [View history](#) [Tools](#) ▾

From Wikipedia, the free encyclopedia

"Dieselgate" and "Emissionsgate" redirect here. For other diesel emissions scandals, see [Diesel emissions scandal](#).

The **Volkswagen emissions scandal**, sometimes known as **Dieselgate**^{[23][24]} or **Emissionsgate**,^{[25][24]} began in September 2015, when the [United States Environmental Protection Agency](#) (EPA) issued a notice of violation of the [Clean Air Act](#) to German automaker [Volkswagen Group](#).^[26] The agency had found that Volkswagen had intentionally programmed [turbocharged direct injection](#) (TDI) [diesel engines](#) to activate their [emissions](#) controls only during laboratory [emissions testing](#), which caused the vehicles' [NO_x](#) output to meet US standards during regulatory testing. However, the vehicles emitted up to 40 times more NO_x in real-world driving.^[27] Volkswagen deployed this software in about 11 million cars worldwide, including 500,000 in the United States, in [model years](#) 2009 through 2015.^{[28][29][30][31]}

Volkswagen emissions scandal



A 2010 Volkswagen Golf TDI displaying "Clean Diesel" at the [Detroit Auto Show](#)

Search engine optimization

🌐 65 languages ▾

Article [Talk](#) [Read](#) [View source](#) [View history](#) [Tools](#) ▾

From Wikipedia, the free encyclopedia

"SEO" redirects here. For other uses, see [Seo](#).

Search engine optimization (SEO) is the process of improving the quality and quantity of [website traffic](#) to a [website](#) or a [web page](#) from [search engines](#).^{[1][2]} SEO targets unpaid traffic (known as "natural" or "[organic](#)" results) rather than direct traffic or [paid traffic](#). Unpaid traffic may originate from different kinds of searches, including [image search](#), [video search](#), [academic search](#),^[3] news search, and industry-specific [vertical search engines](#).

As an [Internet marketing](#) strategy, SEO considers how search engines work, the computer-programmed [algorithms](#) that dictate search engine behavior, what people search for, the actual search terms or [keywords](#) typed into search engines, and which search engines are preferred by their targeted audience. SEO is performed because a website will receive more visitors from a search engine when websites rank higher on the [search engine results page](#) (SERP). These visitors can then potentially be converted into customers.^[4]

| |
|-----------------------------------|
| Part of a series on |
| Internet marketing |
| Search engine optimization |
| Local search engine optimisation |
| Social media marketing |
| Email marketing |
| Referral marketing |
| Content marketing |
| Native advertising |
| Search engine marketing |
| Pay-per-click |
| Cost per impression |
| Search analytics |
| Web analytics |
| Display advertising |
| Ad blocking |
| Contextual advertising |

Errors of assumptions around tools



HOME
ABOUT EuSpRIG
EuSpRIG 2019 ANNUAL CONFERENCE
EuSpRIG 2019 PROGRAMME
REGISTER FOR EuSpRIG 2019 CONFERENCE
BASIC RESEARCH
BEST PRACTICE
HORROR STORIES
REGULATORS' PRESENTATIONS
CONFERENCE ABSTRACTS, PAPERS & INDEXES
CONFERENCE REPORTS & VIDEOS
DELEGATES

CONSTITUTION
HISTORY
OUR TALKS & PRESENTATIONS
QUOTABLE QUOTES
PRESS & WEBSITE
COMMITTEES
DISCUSSION GROUP
TRAINING VIDEOS
HUMOUR
SPONSORS
USEFUL LINKS
CONTACT



EuSpRIG Horror Stories

Spreadsheet mistakes - news stories

Public reports of spreadsheet errors have been sought out on behalf of EuSpRIG by Patrick O'Beirne of Systems Modelling for many years. There are very many reports of spreadsheet related errors and they seem to appear in the global media at a fairly consistent rate.

These stories illustrate common problems that occur with the uncontrolled use of spreadsheets. In many cases, we identify the area of risk involved and then say how we think the problem might have been avoided.

Stories are identified by those who kindly collated and sorted them:

POB: Patrick O'Beirne, Eusprig chair

FH: Fellenne Hermans (winner of the 2011 [David Chadwick student prize](#) and now an assistant professor at Delft University of Technology).

NS: Tie Cheng, a EuSpRIG [committee member](#).

MPC: Mary Pat Campbell, an actuary, trainer, and a member of the [EuSpRIG Discussion group](#).



Identifier: POB2001
Title: Data not controlled, 16000 UK Covid-19 test results lost for a week
Source: <https://www.bbc.co.uk/news/technology-54423988>
Release Date: 08 October 2020
Risk: Lives put at risk because the contact-tracing process had been delayed
Discrepancy: 16,000 test cases in a week

Excel: Why using Microsoft's tool caused Covid-19 results to be lost

"The badly thought-out use of Microsoft's Excel software was the reason nearly 16,000 coronavirus cases went unreported in England. [The labs] filed their [result logs] results in the form of text-based lists - known as CSV files - without issue. PHE had set up an automatic process to pull this data together into Excel templates so that it could then be uploaded to a central system. The problem is that [Public Health England] PHE's own developers picked an old file format to do this - known as XLS. As a consequence, each template could handle only about 65,000 rows of data rather than the one million-plus rows that Excel is actually capable of. And since each test result created several rows of data, in practice it meant that each template was limited to about 1,400 cases. When that total was reached, further cases were simply left off. To handle the problem, PHE is now

<http://www.eusprig.org/horror-stories.htm>



Ads by Google

Stop seeing this ad

Why this ad? ▸

MICROSOFT REPORT SCIENCE

Scientists rename human genes to stop Microsoft Excel from misreading them as dates

99

Sometimes it's easier to rewrite genetics than update Excel

By James Vincent | Aug 6, 2020, 8:44am EDT



Listen to this article



SHARE



Effective, Reliable, Secure

AWS Solutions Implementations help you build faster with vetted architecture.

Get Started

What about when Excel is used as intended?



15k spreadsheets
97M cells
20M formulas

Enron's Spreadsheets and Related Emails: A Dataset and Analysis



Felienne Hermans
Delft University of Technology
Mekelweg 4
2628 CD Delft, the Netherlands
f.f.j.hermans@tudelft.nl

Emerson Murphy-Hill
North Carolina State University
890 Oval Drive
Raleigh, North Carolina, USA
emerson@csc.ncsu.edu



TABLE III
SPREADSHEETS CONTAINING EXCEL ERRORS IN THE ENRON SET

| Error type | Spreadsheets | Formulas | Unique Ones |
|------------|--------------|-----------|-------------|
| #DIV/0! | 580 | 76,656 | 4,779 |
| #N/A | 635 | 948,194 | 6,842 |
| #NAME? | 297 | 33,9365 | 29,422 |
| #NUM! | 52 | 4,087 | 178 |
| #REF! | 931 | 18,3014 | 6824 |
| #VALUE! | 423 | 11,1024 | 1751 |
| Total | 2,205 | 1,662,340 | 49,796 |

24% of spreadsheets with formulas had errors



Learning the irrelevant ML loves the shortcut



Article: Super Bowl 50

Paragraph: "Peyton Manning became the first quarterback ever to lead two different teams to multiple Super Bowls. He is also the oldest quarterback ever to play in a Super Bowl at age 39. The past record was held by John Elway, who led the Broncos to victory in Super Bowl XXXIII at age 38 and is currently Denver's Executive Vice President of Football Operations and General Manager. Quarterback Jeff Dean had a jersey number 37 in Champ Bowl XXXIV."

Question: "What is the name of the quarterback who was 38 in Super Bowl XXXIII?"

Original Prediction: John Elway

Prediction under adversary: Jeff Dean

| Task for DNN | Caption image | Recognise object | Recognise pneumonia | Answer question |
|-----------------|---|---|-----------------------------------|---|
| Problem | Describes green hillside as grazing sheep | Hallucinates teapot if certain patterns are present | Fails on scans from new hospitals | Changes answer if irrelevant information is added |
| Shortcut | Uses background to recognise primary object | Uses features irrecognisable to humans | Looks at hospital token, not lung | Only looks at last sentence and ignores context |

same category for humans but not for DNNs (intended generalisation)

i.i.d.



domain shift
e.g. Wang '18



adversarial examples
Szegedy '13



distortions
e.g. Dodge '19



pose
Alcorn '19



texture
Geirhos '19



background
Beery '18



o.o.d.



same category for DNNs but not for humans (unintended generalisation)



excessive invariance
Jacobsen '19



fooling images
Nguyen '15



natural adversarial examples
Hendrycks '19



texturised images
Brendel '19



■ Tell me a novel joke about the singer Madonna.

■ Why did Madonna study geometry?

Because she wanted to learn how to strike a pose in every angle! 📐 🦋

10 Ways GPT-4 Is Impressive but Still Flawed

OpenAI has upgraded the technology that powers its online chatbot in notable ways. It's more accurate, but it still makes things up.

By Cade Metz and Keith Collins

Cade Metz asked experts to use GPT-4, and Keith Collins visualized the answers that the artificial intelligence generated.

March 14, 2023

**Errors around what we
get to know**

Preregistration and publishing negative results

“Data available upon reasonable request”

...and this is where we put the
non-significant results.

som^{ee}cards
user card








Figure 1: The distribution of more than one million z -values from Medline (1976–2019).

p-values in published research



What explains this distribution?






Figure 1: The distribution of more than one million z -values from Medline (1976–2019).




Figure 1: The distribution of more than one million z -values from Medline (1976–2019).

